

LIVE DEMONSTRATION SCRIPT

HORIZON GRID

Hackathon Demo Guide

Version: 0.2.4

Documentation prepared: 2026-08-21

PREPARED FOR EVALUATION

Table of Contents

HORIZON GRID — Hackathon Demo Guide	3
-------------------------------------	---

HORIZON GRID — Hackathon Demo Guide

Every Signal. One Operational Picture.

A suggested 7–10 minute live-demo sequence, built around what the product actually does today (v0.2.4). Every step below drives a real feature — nothing here is scripted around a mockup or a canned response.

0. Before You Start

- The stack should already be running (`docker compose ... up -d`) against the installed copy, or via the Windows/Linux installer). Confirm the Executive Dashboard loads real numbers (not zeros-because-nothing-ran) by having at least one completed investigation in the database beforehand.
- Have one AI backend configured and live-tested (Ollama if you want a fully offline demo with no API key, or any of the ten cloud backends if you want faster responses).
- Have at least one IOC provider credential configured beyond the no-key providers, so the demo shows real provider diversity rather than only the six providers that need no key at all (crt.sh, CISA KEV, MITRE ATT&CK, WHOIS/RDAP, Spamhaus, Internet Intelligence Collector).

1. The Problem, in One Sentence (30 seconds)

"No single threat-intel source sees everything, so analysts manually cross-reference five or six of them per indicator — and that manual cross-referencing is the first thing that gets skipped when the queue is long."

2. Executive Dashboard First, Not the Search Box (1 minute)

Open the dashboard before running anything new. Point at the seven live KPI tiles (active investigations, critical/high-risk IOC count, open case counts, average threat score, provider health percentage, 30-day AI success rate) and the Provider Health summary widget. Say explicitly: **every number on this screen is a real query result, not a hardcoded demo value** — that's a deliberate design constraint, not an accident, and it's worth saying out loud because judges have seen dashboards that fake this.

3. Run a Live Investigation (2–3 minutes)

Submit a real IOC from the global search bar — a known-benign IP (e.g. `8.8.8.8`) for a clean demo, or a known-malicious sample hash if you want the more dramatic result. Narrate what's happening as it streams in:

- Results arrive per-provider over Server-Sent Events, not all at once — the page doesn't block waiting for the slowest provider.
- Each provider result gets its own short AI summary as it lands, grounded only in that provider's data.
- Once every provider has reported, point out the two things that happen *before* the AI's final verdict: correlation (shared infrastructure, related hashes, MITRE ATT&CK techniques) and the deterministic scoring engine computing `overall_risk_score` / `confidence_score` / `malicious_probability` / severity from that evidence.
- Only then does the AI's consolidated Final Assessment appear — narrating the score it was handed, not inventing one.

4. Prove the AI Isn't the Source of Truth (1–2 minutes)

This is the single most differentiating moment in the demo. Open the Evidence Ledger and one of the verdict-interrogation tools ("Why?", "Challenge This Verdict", "Score Explanation", or "Intelligence Conflicts") and show that every AI claim traces back to a concrete, non-AI-generated fact. Then switch the active AI backend (AI Quickswitch, top of the home page) to a different provider and re-run the same investigation's assessment — the threat score stays identical because it was never computed by the AI in the first place; only the narration changes.

5. Provider Health (1 minute)

Navigate to the full Provider Health page. Show the four rolling windows (1h/24h/7d/30d) and point out the honesty rule that matters most here: a provider with zero real attempts in a window shows **Unknown**, never a false "Healthy" — and a provider correctly returning "nothing found" counts as a healthy outcome, not a failure. This is the page that would catch a Safe-Browsing failure being silently misread as "safe," which is exactly the kind of dashboard lie this design is built to prevent.

6. Security Assessment Toolkit (1–2 minutes, optional but strong)

From a completed investigation, trigger the Security Assessment Toolkit against a target you're explicitly authorized to scan (your own demo host is safest). Show the two required confirmations (retype the target, confirm authorization) before anything runs — call out that this is the one feature in the product that sends real active traffic, unlike every passive provider lookup, and that it's gated accordingly. Show a finding feeding back into the investigation's risk/confidence floor without inflating `malicious_probability` on its own.

7. Administration and RBAC (1 minute)

Open the Admin console. Show the three fixed roles (Admin/Analyst/Viewer) and that permission checks are enforced server-side on every route, not just hidden in the UI — a Viewer can see the dashboard and provider health but cannot run a scan or export. If more than one admin account exists, mention that multi-admin support is real, not a single hardcoded superuser.

8. Export and Close (30 seconds)

Export the investigation you ran (PDF or Markdown) and point out the HORIZON GRID branding, generation timestamp, and Investigation ID baked into the export — every report carries provenance, so a downstream reader can trace it back to a specific run.

Closing Line

"Every number you saw today — the dashboard KPIs, the threat score, the provider health page — was computed live from real evidence, and every AI statement was checkable against that evidence. That auditability, not the AI itself, is the actual product."

If Something Goes Wrong Live

- **A provider is down/rate-limited mid-demo:** that's fine — say so, and point at Provider Health showing the same thing honestly. It's a better demo moment than a fake all-green dashboard would be.
- **The configured AI backend fails:** the platform is designed to distinguish a real AI failure from a correct "no evidence to reason over" skip, with a disclosed fallback either way — narrate that distinction live if it happens rather than treating it as a demo failure.