

HACKATHON SUBMISSION PACKAGE

HORIZON GRID

Hackathon Submission — Every Signal. One Operational Picture.

Version: 0.2.4

Documentation prepared: 2026-08-21

PREPARED FOR EVALUATION

Table of Contents

Executive Summary	3
The Problem and The Solution	5
Feature Inventory	7
The Development Journey	11
Design Evolution	14
Brand Identity	16
Bug History: What Actually Broke, and What Got Fixed	19
The Testing Journey	25
Roadmap and Conclusion	30

Executive Summary

HORIZON GRID is a self-hosted threat-intelligence platform. An analyst gives it one indicator of compromise — an IP address, domain, URL, file hash, or CVE ID — and it queries 18 independent intelligence sources at once, correlates what they say about each other, computes a deterministic threat score from that evidence, and has an AI model explain the result in plain language, with every claim traceable back to the real data behind it. The product’s own tagline states the design goal directly: **Every Signal. One Operational Picture.**




Figure 1 — The Executive Dashboard — every KPI tile is a real query result against the live database, never a hardcoded display value.

What Problem Does It Solve?

No single threat-intelligence source sees everything. A file hash flagged by one antivirus engine might be ignored by another; an IP address on one blocklist might look unremarkable on the next. Getting a trustworthy read on an indicator normally means an analyst manually checking it against several independent sources, one browser tab at a time, and holding all of it in their head well enough to notice when two sources are describing the same underlying fact from different angles. That manual cross-referencing is the real cost of threat triage — and it is exactly the cost that gets skipped when the queue of indicators is long.

What Does HORIZON GRID Do About It?

It collapses that workflow into a single search box. One submission fans out to every applicable provider in parallel; results correlate automatically; a fixed, versioned, non-AI scoring engine turns the combined evidence into a threat score before any AI is involved; and an AI model is then handed that score as a given fact to narrate, not a number it's free to invent. The platform tracks its own AI outcomes honestly — a real success, a correct decision to skip the AI when there's no evidence to reason over, or a disclosed failure with a deterministic fallback — never a fabricated result.

Who Uses It?

A solo analyst can run the whole stack on one Windows or Linux machine behind a guided installer; a small security team can point several browsers at the same install and share cases, a shared watchlist ("the Basket"), and one live operational dashboard. Three fixed roles — Admin, Analyst, Viewer — are enforced on every backend route, not just hidden in the UI, so a broader audience (including a Viewer who should see the operational picture but not run scans or export data) can be given access safely.

What Makes It Different?

Two things a typical dashboard tool doesn't do:

1. **The threat score is computed before the AI ever sees it, not by the AI.** A deterministic scoring engine (versioned, currently `1.0`) turns provider-verdict consensus and correlation-graph evidence into `overall_risk_score`, `confidence_score`, `malicious_probability`, and a severity band using a fixed, documented formula. The AI is handed that number as an input and narrates it — the platform re-validates the AI's own stated verdict against that number before saving anything, so the AI cannot quietly override the math.
2. **Every AI claim is checkable against real evidence, not just plausible.** An Evidence Ledger and a set of verdict-interrogation tools ("Why?", "Challenge This Verdict", "False Positive Check", "Score Explanation", "Intelligence Conflicts") let an analyst check any AI statement against the concrete, non-AI-generated facts the platform actually collected — rather than trusting a well-written paragraph on faith.

Why It Matters

Analysts don't need another tool that produces a confident-sounding paragraph about an indicator — they need one that shows its work. HORIZON GRID's central design bet is that a security tool earns trust by being *auditable*: a fixed score computed the same way every time, an AI that narrates rather than decides, and a health/operational picture (Provider Health, Executive Dashboard) that is never allowed to show a hardcoded or invented number. That bet shaped every architectural decision documented in the rest of this submission.

The Problem and The Solution

The Problem: Fragmented, Manual Threat Triage

Every individual threat-intelligence source is narrow by design — a multi-engine reputation scanner, a malware sandbox, a certificate-transparency search, a government vulnerability catalog, a DNS blocklist — and none of them sees the full picture alone. In practice, investigating one indicator means an analyst opens a reputation lookup site, then a sandbox or sample database, then a certificate search, then a vulnerability database, then a blocklist, pasting the same value into each one and reading each result in its own format. Cross-referencing those results by hand — noticing that an IP tied to a malware family is also tied to a MITRE ATT&CK technique that matches a hash flagged somewhere else — is the actual cost of triage, and it is the first thing that gets skipped when the indicator queue is long.

A second, quieter problem sits underneath the first: once AI entered this workflow, it became tempting to let a language model produce the whole verdict from a prompt. That's fast, but it trades a real problem (manual cross-referencing is slow) for a worse one — an analyst now has to trust a paragraph that might be internally inconsistent, might invent a source that was never checked, and offers no fixed, reproducible number to audit later.

The Solution: One Submission, One Verifiable Operational Picture

HORIZON GRID's actual processing pipeline, in the order it really executes:

1. **IOC submission and classification** — the analyst submits one value; the platform validates and classifies it (IPv4/IPv6, domain, URL, file hash, CVE ID, and the other IOC types the app recognizes).
2. **Parallel provider queries** — every one of the 18 built-in providers relevant to that IOC type is queried at once (a file hash is never sent to a domain-registration lookup), and results stream into the page live over Server-Sent Events as each provider responds, rather than waiting for the slowest one.
3. **Per-provider AI summaries** — as each provider's result arrives, a short AI-written summary of that one result is generated, grounded only in that provider's own data.
4. **Correlation** — once providers have reported in, a correlation engine extracts relationships between what different sources said (shared infrastructure, related hashes, malware families, MITRE ATT&CK techniques, exploited CVEs) and boosts confidence when independent providers corroborate the same fact.
5. **Deterministic threat scoring — before the AI's final verdict, not after.** A fixed, versioned scoring engine computes `overall_risk_score`, `confidence_score`, `malicious_probability`, and a severity band from provider-verdict consensus and correlation evidence, using a documented, weighted formula (see the Threat Scoring guide). This step deliberately happens *before* the AI's consolidated assessment.
6. **AI Final Assessment, conditioned on the score** — one consolidated AI pass reasons over every per-provider summary plus the correlation output, and is handed the already-computed score as a fixed input to narrate. The platform validates that the AI's own stated verdict agrees with that number before saving anything, rather than trusting the AI to have applied the number correctly.
7. **Operational presentation** — the result renders as an investigation page with an Evidence Ledger and verdict-interrogation tools, rolls up into the Executive Dashboard's live KPIs, and can be exported (JSON/Markdown/CSV/PDF) with the platform's branding, a generation timestamp, and an Investigation ID on every export.

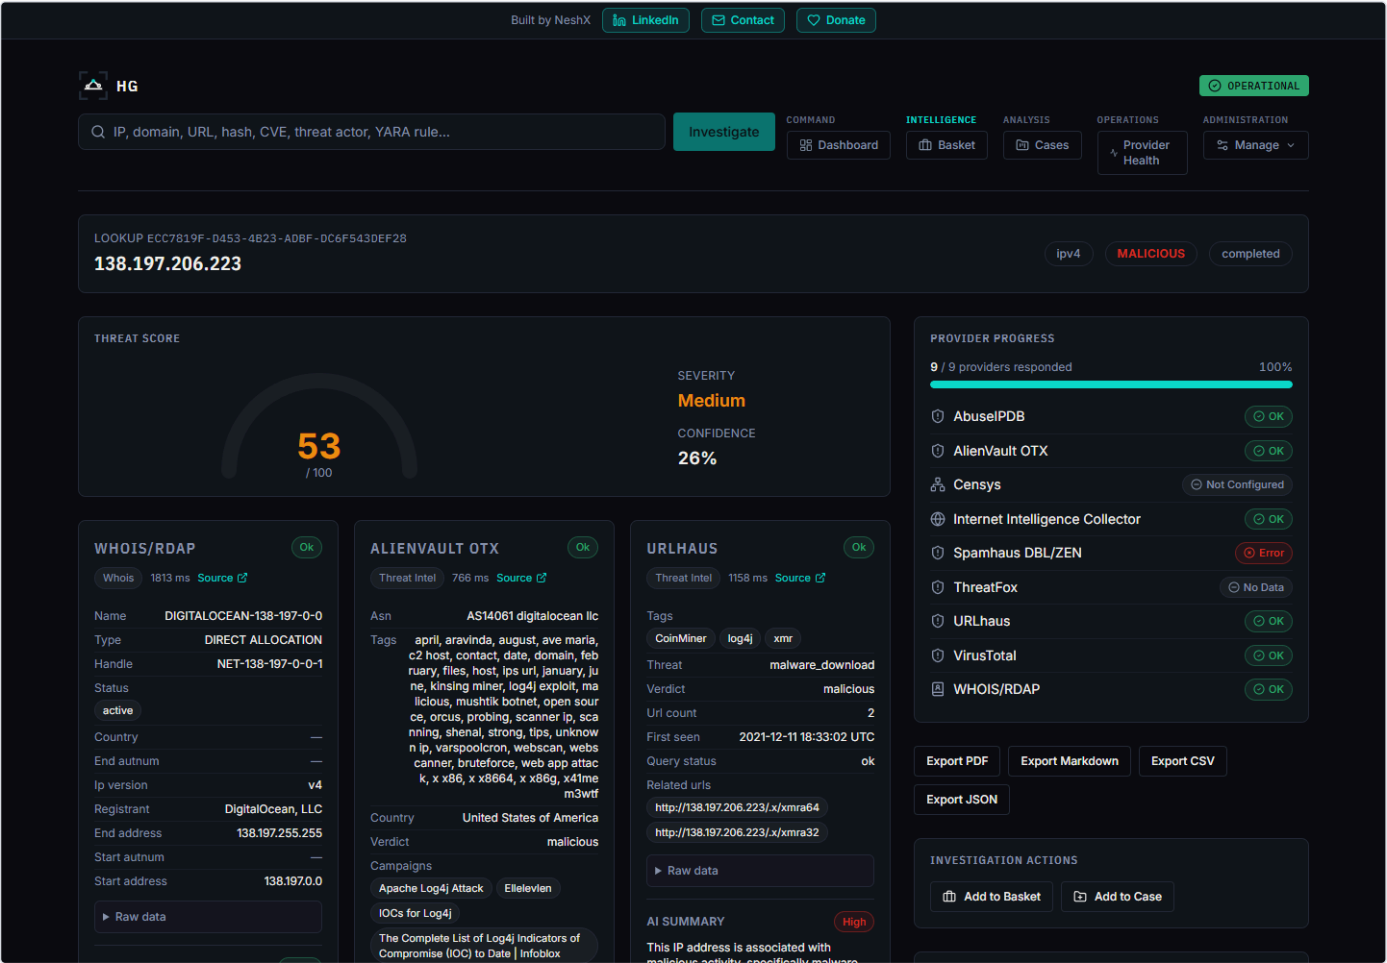


Figure 2 — A completed investigation — provider results, threat score, and severity all visible on one page, each traceable back to the specific source that produced it.

The order in step 5 and step 6 is the load-bearing design decision in this whole pipeline: the AI is never the source of the threat score. It receives a number that was already computed by a fixed formula and is only permitted to explain it — which is what makes the score reproducible and auditable in a way a single AI-generated verdict never can be.

Why This Design, Specifically

An earlier, simpler version of this idea would have let the AI produce `risk_score`, `malicious_probability`, and the verdict directly from a prompt in one call. That's a real, common pattern, and it's exactly the pattern this project moved away from once a deterministic scoring engine was built: prompt wording alone cannot reliably stop a model — particularly a local model — from overriding a "do not fill in, this is provided" instruction. Computing the score first and feeding it into the AI prompt as a fixed input, then validating the AI's verdict against it, closes that gap structurally instead of relying on the AI to behave. It is also the only design that makes an "average threat score" KPI on the Executive Dashboard mean anything consistent over time, since the AI backend an analyst has selected (of eleven interchangeable options) no longer changes what the score is — only how it's explained.

Feature Inventory

Every item below is implemented and verified against current source code (`backend/app/providers/registry.py` , `backend/app/core/runtime_config.py` , `backend/app/scoring/engine.py` , `backend/app/models/user.py` , `backend/app/security_assessment/`) — nothing here is planned or aspirational. Deep-dive detail on each area lives in the companion technical documentation set (Provider Guide, AI Guide, Threat Scoring guide, Admin Guide, Security guide) referenced at the end of this submission.

IOC Investigation

The core workflow: submit one indicator (IP, domain, URL, file hash, CVE, and 27 other recognized `IOCType` values — 32 in total), and the platform validates, classifies, and fans the value out to every registered provider that supports that type. Results stream in over Server-Sent Events rather than blocking on the slowest provider.

Provider Ecosystem — 18 Providers

`VirusTotal` , `AbuseIPDB` , `AlienVault OTX` , `URLhaus` , `ThreatFox` , `MalwareBazaar` , `crt.sh` , `NIST NVD` , `CISA KEV` , `MITRE ATT&CK` , `WHOIS/RDAP` , `urlscan.io` , `Google Safe Browsing` , `Hybrid Analysis` , `Spamhaus` , `PhishTank` , `Censys` , and the platform's own **Internet Intelligence Collector** (a live OSINT crawler) — spanning multi-engine reputation scanning, community abuse/malware databases, certificate-transparency search, government vulnerability catalogs, sandbox behavior, DNS blocklists, phishing databases, and internet-wide host scanning. Several (`crtsh` , `cisa_kev` , `mitre_attack` , `whois_rdap` , and the stub-implemented `spamhaus` / `phishtank`) need no API key at all and are active from first boot.

AI System — 11 Interchangeable Backends

`Ollama` (local, free, no key), `Anthropic` , `AWS Bedrock` , `Google Gemini` , `Groq` , `OpenAI` , `Kimi` (Moonshot AI), `DeepSeek` , `xAI` (Grok), `Mistral AI` , and `OpenRouter` — all exposed through one identical `call_claude_json()` interface, switchable at runtime with no restart. Two distinct AI passes run per investigation: a grounded per-provider summary as each result arrives, and one consolidated Final Assessment once every provider has reported, conditioned on the deterministic threat score (see below). Every AI generation is tracked as one of three honest outcomes — real success, correct no-evidence skip, or disclosed failure — never a fabricated result.

Deterministic Threat Scoring

A fixed, versioned (`SCORING_ENGINE_VERSION = "1.0"`) scoring engine computes `overall_risk_score` , `confidence_score` , `malicious_probability` , and a severity band from provider-verdict consensus and correlation-graph evidence — **before** the AI's final assessment runs, not after. The AI receives the score as a fixed input and is validated against it, so it narrates rather than decides. A dedicated corroboration-scaling mechanism prevents one account listing several tags across free-text fields from single-handedly flooding the score.

Correlation and Evidence

A pure, deterministic correlation engine extracts typed relationship edges (shared infrastructure, related hashes, malware families, MITRE ATT&CK techniques, exploited CVEs) from normalized provider fields and boosts confidence when independent providers corroborate the same fact. A non-AI-generated Evidence Ledger, built directly from provider results and correlation edges, backs a set of verdict-interrogation tools (WHY malicious, What Is This, Provider Disagreement, False-Positive Assessment, Challenge/red-team, Smart Next Actions, Intelligence Gaps, Score Explanation) plus Copilot Q&A, IOC comparison, and hunting-query/detection-rule generation.

Executive Dashboard

Live KPI tiles — active investigations, critical/high-risk IOC count, open case counts, average threat score, provider health percentage, AI success rate — computed from real database aggregates, never hardcoded, plus an AI-generated narrative that is honestly labeled as AI-generated or template-fallback.




Figure 3 — The Final Assessment panel's Risk & Verdict tab — overall risk score, confidence score, severity, and malicious probability, all computed before the AI narrates them.

Provider Health

A dedicated page tracking every registered provider's real status (Healthy/Degraded/Down/Unknown) across four rolling windows (1h/24h/7d/30d), built from the same `ProviderResultRecord` data every real investigation already writes. A provider with zero attempts in a window shows Unknown, never a false Healthy.

Security Assessment Toolkit (Active Port Scanner)

A genuinely active check — Nmap port/service scan, DNS record lookup, TLS certificate inspection, HTTP security-header check — against a target the analyst explicitly confirms authorization for (target retype + authorization checkbox), with real cancellation support (a live task can be stopped mid-scan, not just flagged in the database). Findings feed the investigation's risk/confidence floor without inflating `malicious_probability` on their own.

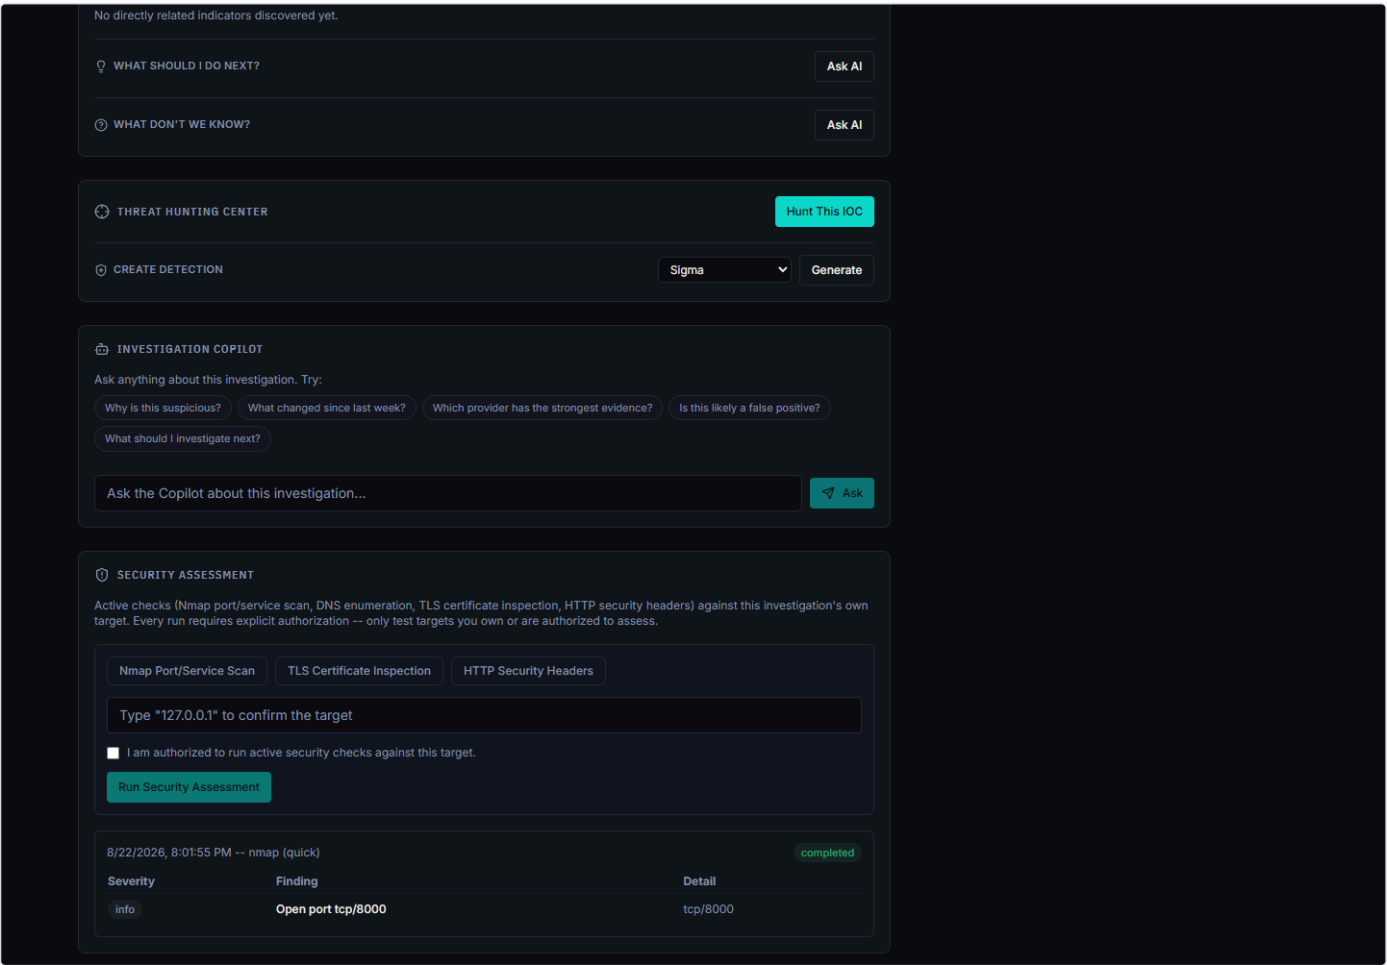


Figure 4 — A completed Nmap "quick" scan against an authorized target, showing the required target-confirmation/authorization controls and a real finding.

Case Management and IOC Basket

Group indicators into a case with analyst notes as an incident develops, or keep a running personal watchlist of indicators worth revisiting.

Administration and RBAC

Three fixed roles — Admin, Analyst, Viewer — enforced server-side on every route via a `require_permission()` dependency, not just hidden in the UI. Genuine multi-administrator support (any admin can create further admin accounts from the console; the platform does not hardcode a single superuser). A full audit log records configuration and security-relevant changes.

Exporting Findings

JSON, Markdown, CSV, and PDF export, gated behind an export-specific permission distinct from ordinary read access. CSV and PDF paths carry tested defenses against formula-injection and markup-injection abuse. Every export carries the platform's branding, a generation timestamp, and an Investigation ID.




Figure 5 — The export panel on a completed investigation — JSON, Markdown, CSV, and PDF, each carrying platform branding and an Investigation ID.

Deployment

A Windows Inno Setup installer with a guided WinForms Setup Wizard (administrator account, AI backend, provider credentials with live "Test" buttons, network ports, install), and a Linux `.deb` package with an equivalent setup flow — both running the actual application as Docker Compose, plus a parallel Kubernetes/Kustomize deployment alternative for either platform.

The Development Journey

Before Version Control

The product already existed as a working prototype before formal version control began — the earliest QA documents in this repository (`BACKEND_TEST_ROOT_CAUSE_REPORT.md` , `API_CONFIGURATION_FIX_REPORT.md` , `ADMIN_RBAC_QA_REPORT.md` , `RUNTIME_PROVIDER_IMPLEMENTATION_REPORT.md`) describe a system with a working Setup Wizard, a database-backed RBAC console, and a runtime-mutable AI/provider configuration layer, none of which are attributed to any commit in this repository's history. That earlier phase is not independently documented here, and this narrative does not attempt to reconstruct it — verifiable history begins at git's own root commit, `9bee96c` , "Initial commit: HORIZON GRID threat intelligence platform," recorded 2026-08-18 21:37.

The Initial Commit

By the time `9bee96c` was recorded, the platform was already substantial, not a skeleton. Reading the committed tree directly: `README.md` already opens with the name "HORIZON GRID" and the tagline "Every Signal. One Operational Picture."; the deterministic scoring engine (`backend/app/scoring/engine.py`) and the dashboard aggregation module (`backend/app/core/dashboard.py`) both trace back to this exact commit with `git log --follow` — they were not added later. The frontend's own `package.json` , however, still carried the internal name `"ioc-intel-platform-frontend"` at this point, a naming inconsistency that outlived several later releases. This single commit is the entire starting baseline for every subsequent phase described below; there is no earlier commit to compare it against.

The Adversarial Audit

The next five commits, all recorded within about ninety minutes of each other late on 2026-08-18, form a self-escalating hardening cycle rather than a single fix.

`c222d2f` ("Add GitHub repo scaffolding: README, SECURITY.md, CI, templates; fix real issues found along the way") added the GitHub-facing scaffolding — issue templates, a PR template, CODEOWNERS, and three CI workflows — grounded in a fresh source-code audit that counted 17 IOC providers plus one OSINT crawler and 6 AI backends at that point. The same commit closed real dependency advisories (`python-jose` 3.3.0→3.4.0 , `cryptography` 43.0.1→49.0.0 , `python-multipart` 0.0.9→0.0.31) and fixed a genuinely broken `frontend/.eslintrc.json` — Next.js's linter had no config at all and was silently unusable, and adding one surfaced two real `react/no-unescaped-entities` errors in `NetworkAccessPanel.tsx` .

`54362b2` is a CI-only bugfix: a new Docker-based integration test job ran the test suite against a genuinely empty database for the first time (rather than the developer's own persistent dev database, which already had an admin account), and `test_role_downgrade_immediately_revokes_the_old_roles_access` failed because promoting the test's lone user to admin made them the sole active admin — demoting them then tripped the platform's own last-admin-protection guard instead of exercising the intended role-revocation path. Fixed by creating a companion admin for the duration of the test.

`f305430` ("Adversarial audit fixes: undocumented CVEs, leaked test credential, stale doc counts") is explicitly described in its own commit message as a second, skeptical re-audit that assumed nothing from the prior pass was correct. It found that the `cryptography` version just bumped to 49.0.0 had its own new high-severity advisory (CVE-2026-69247, a PKCS7 Bleichenbacher oracle), requiring a further bump to 50.0.0, and that `pyasn1` — pulled in transitively via `python-jose` and previously untracked — carried four high-severity DoS advisories on the resolved 0.4.8. It scrubbed a real plaintext test admin/analyst password out of `BUG_Triage.md` , `ENTERPRISE_QA_PHASE1_FINDINGS.md` , and `test-provider-pipeline.sh` (the script now reads `TEST_ADMIN_EMAIL` / `PASSWORD` from the environment instead of hardcoding a credential), and found that `.env.example`'s `JWT_SECRET_KEY` placeholder didn't match `config.py`'s own `Settings` default — a fresh install copying `.env.example` verbatim would boot with a well-known, guessable JWT signing key. It also corrected a stale claim that the scoring engine's test suite had 37 tests; direct count showed 22.

`99db66e` is the direct consequence of `f305430`'s own `pyasn1` pin: CI caught, on push, that `python-jose==3.4.0` declares `pyasn1<0.5.0` , directly conflicting with the `pyasn1==0.6.4` pin just added — a real `pip install -r requirements.txt` hard-fails with `ResolutionImpossible` . The commit message is explicit that the prior local verification only ran an incremental `pip install --upgrade` on an already-resolved virtualenv, which doesn't re-validate the full dependency graph the way a fresh install does — "exactly the gap a real CI run exists to catch." Fixed by bumping `python-jose` to 3.5.0, which itself relaxed its own constraint to `pyasn1>=0.5.0` .

`0ba0b1c` root-caused an intermittent "works after reinstall" test flakiness to two divergent, undocumented local virtualenvs (`backend/.venv` , `backend/.venv_test`) with no record of which was canonical — `.venv` was missing `cryptography` entirely and had drifted to unpinned latest-on-PyPI versions of `fastapi` / `pytest` / `starlette` . A second, independent root cause was identified: the pinned `asynccpg==0.29.0` has no prebuilt wheel for Python 3.13 on Windows, which had forced whoever set up each `venv` to work around it differently. The fix bumped `asynccpg` to 0.30.0, deleted the broken `.venv` , and declared `.venv_test` the one canonical dev environment going forward, documented in a new `BACKEND_TEST_ROOT_CAUSE_REPORT.md` . That same report discloses, rather than hides, one pre-existing and still-unfixed issue: cross-file event-loop teardown flakiness in the integration test suite, which does not affect the unit suite CI actually gates on.

AI Provider Expansion — v0.2.0

`e6539ad` ("feat: expand AI provider ecosystem 6->11 backends, v0.2.0"), recorded 2026-08-19 00:40, is the commit that formally bumped the version from 0.1.0 to 0.2.0 — across `windows/installer.iss` , `linux/debian/control` , `backend/app/main.py` , and `frontend/package.json` . Its actual code change, confirmed by both the commit message and `backend/app/core/runtime_config.py` line 29 (`AI_BACKENDS = ["ollama", "anthropic", "bedrock", "gemini", "groq", "openai", "kimi", "deepseek", "xai", "mistral", "openrouter"]`), was adding five new AI backends — Kimi/Moonshot, DeepSeek, xAI/Grok, Mistral AI, and OpenRouter — to the six that already existed (Ollama, Anthropic, Bedrock, Gemini, Groq, OpenAI), following the existing forced-tool-calling client pattern established by `groq_client.py` / `openai_client.py` . Each backend's real base URL and auth format were checked live against the provider's own current documentation rather than assumed: DeepSeek's base URL has no `/v1` segment, unlike every other backend; xAI's error body is a flat `{"code", "error"}` shape rather than the nested shape most backends use; Kimi's default model was deliberately set to `kimi-k2.5` rather than one of Moonshot's "thinking" models, which reject a forced tool call with an HTTP 400. The commit deliberately did not add Fireworks AI (its model discovery needs a non-key account ID, a real pattern deviation from every other backend) or Cerebras (a 2-model catalog with no documented error-body schema) — a curated exclusion, not an oversight. The same commit fixed a real, previously-flagged installer packaging bug: `installer.iss` had never excluded `backend/.venv` / `.venv_test` from the bundled files, so every prior Windows installer silently shipped the local dev virtualenv, dropping the installer from roughly 69 MB to roughly 8 MB with no functional change. `637b2ef` followed six minutes later, adding `HORIZON_GRID_AI_PROVIDER_EXPANSION_QA_REPORT.md` to document the new backends' live connection tests, model-discovery fallback lists, and a full regression run reported as 313 tests passed (up from 282).

It's worth being precise about scope here: `CHANGELOG.md`'s own `[0.2.0]` entry additionally lists the deterministic scoring engine, the executive dashboard, two new IOC providers (urlscan.io, Google Safe Browsing), and the visible rebrand to "HORIZON GRID" as part of this release. Checking those claims against git history directly shows the scoring engine and dashboard modules were already present at the Initial commit (`git log --follow` traces both files back to `9bee96c` , not to `e6539ad`), and `CHANGELOG.md` itself states elsewhere that "pre-0.2.0 work (initial platform build, GitHub scaffolding, early bug fixes) is captured in the git history itself... rather than a version-numbered changelog entry, since HORIZON GRID's first numbered release is 0.2.0." In other words: v0.2.0 is the point where formal version-numbered release documentation began, and its changelog entry retroactively describes the full state of a build whose scoring engine, dashboard, and README-level branding already existed beforehand; the AI backend expansion (6→11) and the installer virtualenv fix are the specific changes this pair of commits actually shipped.

Port-Scanner Cancellation Hardening — v0.2.1 and v0.2.2

Six commits across 2026-08-19 evening and 2026-08-20 early morning cover the Security Assessment Toolkit's port scanner.

[da0d692](#) ("fix: add real cancellation to the Security Assessment port scanner") describes a forensic audit of the full pipeline — frontend, API, validation, the `nmap` subprocess, XML parsing, UI — that found every stage already working correctly against real local targets, with one real gap: once started, a scan could never be stopped short of restarting the backend. It added `POST /api/v1/security-assessment/runs/{run_id}/cancel` and a "Cancel Scan" button, backed by a run-id-keyed dict of live `asyncio.Task` objects so a cancel request can reach an in-flight task and actually kill the OS-level `nmap` process, not just flip a database flag. Two edge cases are called out explicitly in the commit message: `asyncio.CancelledError` does not inherit from `Exception` in Python 3.8+, so it needed its own dedicated `except` branch in both the run orchestrator and the nmap tool's subprocess handling; and a run orphaned by a backend restart (no live task in the current process) still resolves to `CANCELLED` via a direct, rowcount-guarded database write. A migration bug was caught by a genuinely failing test before release — the new `CANCELLED` enum value was first added lowercase, not matching the column's existing uppercase convention — and fixed before the commit. This shipped as v0.2.1.

[5b0edb9](#) fixed a concurrency race found by adversarial testing: two concurrent cancel requests for the same run could each write their own `security_assessment.run_cancelled` audit row, making one real cancellation look like two in the audit trail. The fix makes the run's own background task the single source of truth for that audit event on the live-task path, while the backend-restart-orphan path still writes its own record, gated on its direct database update having actually changed a row. [1c6ced9](#) documented this work in `HORIZON_GRID_PORT_SCANNING_QA_REPORT.md` with a verdict of "PORT SCANNING — RELEASE READY."

That verdict was then independently re-checked rather than taken on faith. [9166e75](#) ("fix: 7 real bugs found by independent re-verification of the cancellation fix") describes nine parallel reviewers re-verifying the v0.2.1 fix from scratch — git archaeology, live browser automation, an adversarial security pass, real concurrency races, and full-application regression — and finding four new defects in the scanner itself plus three unrelated ones surfaced during regression: IPv6 scans never actually probed the target because nmap needs an explicit `-6` flag for an IPv6 literal, and without it silently exits 0 having scanned 0 hosts, reported identically to a real clean scan; an unknown or mistyped scan-profile id was silently accepted and reported as a clean scan with no error anywhere; a malformed CIDR value crashed the run endpoint with a raw HTTP 500; and a run's own successful completion could silently overwrite a status another writer (such as the orphan-recovery sweep) had already finalized. The three unrelated findings were a Spamhaus provider misreading a DNS query-rejection as a positive "malicious" verdict, a confusing generic 401 instead of "Account disabled" on the losing side of a concurrent last-admin-protection race, and an AI-generated assessment that could cite specific findings in prose while leaving its structured supporting-evidence list empty. All seven were fixed, covered by new tests (365 passed, up from 350), and the version was bumped to v0.2.2. [2ef2ad1](#) is a small CI-only follow-up: the new IPv6 argv tests failed on CI's bare "unit" job, which has no `nmap` binary installed, because the tests' mock didn't also cover `is_available()` — fixed by mocking that too. [2e42a14](#) documented the nine-reviewer pass as an addendum to the existing QA report, keeping the same "PORT SCANNING — RELEASE READY" verdict but on stronger evidence.

Reading the resulting code directly (`backend/app/core/security_assessment.py`, `backend/app/security_assessment/nmap_tool.py`) confirms these fixes are actually present: `cancel_run()` and a matching `except asyncio.CancelledError` branch in `_execute_run()` exist with the described audit-dedup gating, `nmap_tool.py`'s `run()` conditionally adds `-6` for IPv6 targets, and `start_run()` validates the profile id and catches CIDR-parsing `ValueError`s before any run row is created.

Mission-Critical Reliability Hardening — v0.2.3

Fourteen commits on 2026-08-20 form the largest single phase in this history: a dedicated review for "unattended remote deployment" — installing the platform once at a remote site with no developer access afterward.

[5491332](#) is the core hardening commit. It added restart policies, log rotation, and memory limits on all 8 Docker Compose services alongside a new dependency-aware `GET /health/detailed` endpoint (checking Postgres and Redis) next to the existing dependency-free `/health` that CI relies on; boot-time auto-start plus a 5-minute health watchdog on both platforms (Windows Scheduled Tasks, Linux systemd timers); and found, in the process, that Linux's `horizon-grid.service` unit file was correct but had never actually been `systemctl enable`d anywhere in the codebase — a working install would silently not survive a host reboot. It wired the existing SSRF guard (`assert_safe_outbound_url`) into the real Ollama AI-call path, including the `.env`-only fallback singleton used by any wizard-driven install that never touches the runtime-config web UI, and into the config-save path rather than only the Test-Connection convenience endpoints. It fixed a dead-retry-code bug where `BaseProvider.run()` was silently catching the exact exception types the orchestrator's retry loop was meant to see, meaning `provider_max_retries` had no real effect for providers that didn't handle their own network errors. It added fixed-window login brute-force rate limiting keyed by attempted email, a concurrency cap on security-assessment scan execution, and CIDR-size-proportional nmap timeout scaling. Separately, it found that `check-prerequisites.sh`'s `add_check()` function interpolated bash's lowercase `true / false` into generated Python source as bare identifiers — Python doesn't recognize those as booleans, so every call silently failed and fell back to a previous value, leaving the script's JSON `checks` array permanently empty (console output and exit code were computed separately and were unaffected, which is why this had gone unnoticed).

[057404e](#) added real database restore scripts on both platforms — previously only backup scripts existed, and the one documented manual restore procedure was self-contradictory (it told an operator to stop the whole platform, then restore into "the running Postgres container," but stopping the platform stops Postgres too). The new scripts stop only the services that write to the database, leave Postgres running, and require a typed confirmation unless `--force / -Force` is passed; this was live-verified end to end (backed up, inserted a marker row, restored, confirmed the marker was gone and the platform came back up correctly). The same commit gave Redis a persistent named volume — confirmed live before the fix that a value set, then a container restart, silently lost it — and fixed `Test-BackendHealth / hg_test_backend_health` (and its frontend equivalent) hardcoding port 8000/3000 regardless of a customized `HOST_PORT_BACKEND / HOST_PORT_FRONTEND`, meaning every unattended caller (the watchdog, service status, restart, open-platform) had been checking the wrong port forever on any non-default install.

[20fc3cc](#) fixed a frontend gap where `FinalAssessmentPanel` rendered the identical badge for `ai_outcome="skipped_no_evidence"` (a correct decision) and `ai_outcome="failed"` (a genuine AI generation failure) — the backend already tracked and sent `ai_outcome` in every response, but the frontend's `FinalAssessment` type simply never included the field, so nothing read it. The same commit added a global 10 MB request-body-size limit (checked against `Content-Length` before the body is read) and `max_length` constraints on several previously-unbounded free-text fields (`LookupCreateRequest.value`, case description/note-body fields), none of which had any limit before. [5052813](#) closed an exposure where `/docs`, `/redoc`, and the raw OpenAPI schema were reachable unconditionally with no environment gate, exposing a full map of every route and request/response shape to anyone who could reach the API — gated behind `settings.environment != "production"` and wired into `docker-compose.prod.yml`.

[6aaffb8](#) added a new Mission-Critical Operations Manual and corrected two stale claims found in the existing Operations Guide: its container table had claimed only `celery_worker / celery_beat` auto-restart (true before this session, wrong afterward — all 8 services do), and its health-check section described an unspecified "backend health endpoint" without noting it now specifically hits `/health/detailed`. [997f6b8](#) bumped the version to 0.2.3 and rebuilt both installers, installing the Linux `.deb` as a real upgrade over an existing v0.2.1 instance inside a WSL test harness — the reconfigure wizard ran end to end (prerequisite gate, pre-upgrade backup, port auto-remap around still-running old containers, new `.env`, containers healthy, all three systemd units enabled) — while disclosing that a live elevated Windows install was not performed in this run, because it requires an interactive UAC prompt the autonomous session had no way to satisfy; the Windows installer's contents were verified directly from the Inno Setup build log instead. [2702ca0](#) added the certification report itself, with a final verdict of "MISSION-CRITICAL READY WITH DOCUMENTED LIMITATIONS," covering 17 real gaps found and fixed (2 of them self-discovered during the review rather than flagged by the initial assessment).

Two smaller fixes followed the same day. [d7f13ae](#) fixed a real crash: the relationship graph's underlying `react-force-graph-2d` /d3-force simulation mutates its input data in place once rendering starts, rewriting `x / y / vx / vy / index` onto node objects and replacing each link's string `source / target` with direct node-object references — since `RelationshipGraph` fed the same memoized `graphData` object to both the force graph and the list view, switching to "View as list" after the graph had rendered once crashed the whole page with React error #31 ("Objects are not valid as a React child"). The fix gives the force graph its own shallow-cloned copy of the data. [f2b61f9](#) corrected the certification report's soak-test section, which had been mistakenly reported as "still running" — the real 3-hour, 171-cycle soak test had actually completed, with memory flat across the run and zero spontaneous container restarts.

The final four commits of this phase are documentation reconciliation: `20e5e22` found and corrected real drift — two chapters still described the wizard polling plain `/health` instead of `/health/detailed`, six files claimed no rate limiting existed on `/auth/login` (a real limiter had just been added), and provider/AI-backend counts across a dozen files were stuck at pre-v0.2.0 numbers (16 providers, 5 AI backends) despite urlscan.io, Google Safe Browsing, and six more AI backends never having been folded into the technical reference chapters. `be41309` rebuilt `CHANGELOG.md` from scratch (not appended) covering v0.2.0 through v0.2.3, along with a version audit, release notes, and test-evidence documents. `e27bde5` regenerated every PDF/DOCX and recaptured 19 screenshots (18 refreshed, 1 new) against the live v0.2.3 build, catching a real bug in the process: the documentation build script's cover-page template had "0.1.0" hardcoded as the version stamp on every previously published PDF, regardless of the actual release — now driven from a single version constant. Two final follow-ups closed out the phase after a real CI failure surfaced post-push: `40ee62b` fixed two Ollama SSRF regression tests that depended on resolving `host.docker.internal` via genuine DNS, which only resolves inside a Docker Desktop network and broke GitHub Actions' bare Ubuntu runner — fixed by mocking the DNS-resolution step to a fixed RFC 5737 test address rather than skipping the check — and `f6be98d` recorded that failure and its fix in the test evidence and changelog, consistent with the release's own stated rule against freezing documentation prematurely.

The threat_assessment Validator Fix

`beb9243` ("fix: reject a bare verdict word in threat_assessment that contradicts final_verdict"), recorded 2026-08-21 04:03, was found live during a full functional-testing pass against a freshly installed instance: investigating `8.8.8.8` with Ollama running `llama3.2:3b` returned `threat_assessment="malicious"` — a bare one-word label, not a sentence — in the same response as a correctly computed `final_verdict="benign"` and `malicious_probability=0.0`. The platform's existing `_verdict_must_agree_with_risk` validator (`backend/app/ai/schemas.py`) only checked `final_verdict` against the numeric risk fields; nothing checked `threat_assessment`'s own text content, and `FinalAssessmentPanel.tsx` renders the two as separately labeled tabs ("Threat Assessment" / "Risk & Verdict"), so an analyst reading either tab alone would see the platform's opposite conclusion.

The fix, read directly from `backend/app/ai/schemas.py` lines 283–309, is a third `model_validator(mode="after")`, `_threat_assessment_bare_verdict_must_agree_with_final_verdict`, deliberately narrow: it only fires when `threat_assessment`, stripped and lowercased, is an exact match for one of `Verdict`'s own string values, and that value's malicious/benign category contradicts `final_verdict`'s. An ordinary narrative sentence that merely contains the word "malicious" mid-sentence never matches this and is unaffected — confirmed by a dedicated regression test asserting exactly that. The validator reuses the same retry-on-`ValidationError` mechanism already in `generate_final_assessment()` rather than introducing new failure-handling logic.

`399c27f` documented this alongside two unrelated P0 installer bugs found during the same interim testing pass — a stale shipped Windows `.exe` missing `docker-compose.yml`, and a leftover verification-harness registry entry silently redirecting installs to a fake path — in `FULL_FUNCTIONAL_TEST_REPORT.md`, explicitly marked "interim, not final," with roughly a third of a 33-phase testing mission covered and the remainder explicitly listed as not yet done. `108d6fa` recorded all three bugs in `CHANGELOG.md`'s v0.2.3 entry along with two new regression tests for the `threat_assessment` / `final_verdict` contradiction — one asserting the exact reproduced payload is rejected, one asserting an ordinary sentence containing the word "malicious" is not.

Final Visual-Identity and Branding Pass — v0.2.4

The last five commits, on 2026-08-21, form a deliberately scoped, presentation-only pass — no backend logic, database schema, authentication, IOC processing, AI provider logic, provider API logic, port-scanner logic, or admin permission was touched, verified by a full backend regression run and a clean frontend `tsc --noEmit` typecheck both before and after.

`c927bf3` ("feat: original visual identity spine (Datum Signal) for the branding pass") is the foundation commit, made because the product name and tagline were, in the commit's own words, "barely visible anywhere in the running application." It introduced a new HSL color palette (a CRT-phosphor teal-cyan primary deliberately pulled off the generic-SaaS-blue hue most dark-mode dashboards default to, against a near-black shell with a warm off-white foreground); an original, wholly geometric logo mark (`Logo.tsx` — a horizon line with two open "signal" nodes converging on one filled "detection" node, chosen specifically because it depicts the tagline rather than merely looking tactical, with no insignia, seal, shield, or existing company/military symbol); a six-state operational status language (`StatusBadge.tsx` : OPERATIONAL/DEGRADED/WARNING/CRITICAL/OFFLINE/UNKNOWN, each state pairing a distinct hue, fill density, border style, and icon — never color alone); a `BrandHeader.tsx` combining the mark with a live SYSTEM STATUS badge polled from the real `GET /health/detailed`; a three-typeface system (IBM Plex Sans Condensed for headings/wordmark/labels, IBM Plex Mono with tabular numerals for raw technical values, body text left unchanged); a radius-contrast rule (soft outer panels, sharp interactive controls) and flat hairline-bordered cards replacing drop-shadows, both applied once at the shared component level; and a new, additive `/about` page.

`a098765` applied that shared header/nav to the ten pages that compose it individually, swapping each page's own `<TopSearchBar /><WorkspaceNav />` pair for `<BrandHeader />` and restyling section headings — explicitly presentation-only, with no API call, state variable, prop signature, or business logic touched on any of the ten files, verified with a clean typecheck after every edit. `f5fac81` extended the branding to every exported report format: PDF exports gained a real header/footer with page numbering via reportlab's canvas callbacks, CSV exports gained a platform/investigation-ID field pair, and the client-built Markdown export gained a matching header — verified live by running a real investigation, exporting the PDF, and visually confirming the header/footer rendered correctly, which caught and fixed a text-overlap spacing bug the unit tests alone hadn't caught. `ce55ec8` bumped the version to 0.2.4.

`1dba60d` closed out the release by syncing `CHANGELOG.md`, the release notes, and the version audit, and — consistent with the documentation-reconciliation pattern seen throughout this history — found six further stale claims while cross-referencing the doc set: three files still said "five supported AI backends" (should be eleven), one said the Linux wizard only offers five AI backends when its own source lists all eleven, and one still said "16 registry providers" in a figure caption whose surrounding prose had already been corrected to 18. Ten canonical screenshot filenames were overwritten with fresh captures against the live v0.2.4 build, and all 25 PDFs/DOCX were rebuilt with the corrected version stamp.

Where This Leaves the Project

Across these 36 commits — from the Initial commit on 2026-08-18 through the branding sync on 2026-08-21 — the recurring pattern is not a single clean build but a series of adversarial re-checks, each one assuming the previous pass might be wrong and finding something real when it looked again: a second audit catching a new CVE the first audit's own fix introduced, nine independent reviewers re-verifying a "release ready" scanner fix and finding seven more bugs, a live investigation of `8.8.8.8` catching an AI self-contradiction no existing validator checked for. The most recent development-cycle document in this repository, `FULL_FUNCTIONAL_TEST_REPORT.md` (2026-08-21, cited above), states its own verdict as interim rather than final, with roughly two-thirds of its planned 33-phase testing mission explicitly marked as not yet run — a disclosed gap, not a claimed one, consistent with how every phase in this history has been documented.

Design Evolution

Where the UI Started

Through the earliest hardening phases (the adversarial audit, the AI provider expansion, the port-scanner cancellation work, and the mission-critical reliability pass — all documented in the Development Journey chapter), the product's UI was functional but generic: a bare search bar and a `WorkspaceNav` top-nav pair repeated at the top of every page, default-case headings, drop-shadowed cards, and a color palette close to the generic SaaS blue most dashboard tools default to. The product name and tagline were barely visible anywhere in the running application — a real gap once the underlying product (17 providers at the time, 11 AI backends, a deterministic scoring engine, an Executive Dashboard, a Provider Health page) had become substantial enough to deserve a real visual identity.

The Branding Pass (v0.2.4)

The last development phase before this submission was a dedicated, presentation-only branding and UI/UX pass, scoped deliberately not to touch backend logic, database schema, authentication, IOC processing, AI provider logic, provider API logic, port-scanner logic, or admin permissions — verified by a full backend regression run (383 passed, 39 skipped, zero regressions) and a clean frontend typecheck both before and after the change.

Three independent visual-identity proposals were generated and judged against each other before implementation began, rather than committing to the first idea that looked reasonable — the same "generate several candidates, synthesize the strongest one" discipline used elsewhere in this project's engineering process. The proposal that won, internally named **"Datum Signal,"** is documented in full in the Brand Identity chapter.

Design Philosophy

The brief for this pass was explicit about the tone to hit and the tone to avoid: the product needed to read as **military-inspired, rugged, operational, mission-ready, precise, technical, and high-trust** — the seriousness of professional aerospace/defense operational software — without literally copying any real defense contractor, military organization, or government agency's branding, insignia, seal, or trademark, and without sliding into a "cyberpunk movie interface" of neon glow and fake-terminal effects. Every element of the final identity was built to satisfy that brief with an *original* mark, palette, and status language rather than borrowed or generic ones.

What Changed, Concretely

- **Color.** A CRT-phosphor teal-cyan primary color was chosen specifically to move away from generic-SaaS blue, paired with a warm off-white foreground against a near-black "anodized steel panel" background.
- **Status language.** A six-state operational vocabulary (OPERATIONAL / DEGRADED / WARNING / CRITICAL / OFFLINE / UNKNOWN) was introduced where every state pairs a distinct hue, fill density, border style, *and* icon — deliberately never color alone, so the product remains legible to a colorblind viewer or on a grayscale printout.
- **Shape language.** A "radius-contrast" rule: outer panels and cards keep a softer corner radius, while every interactive or status element (buttons, inputs, badges, chips) tightens to a sharper radius — soft containers visibly holding precise controls, read as an "operational instrument" without relying on any single color choice to sell that feeling.
- **Elevation.** Drop-shadowed cards were replaced app-wide with a flat fill bounded by a hairline border, so panels read as machined layers rather than floating the way a generic Bootstrap-style dashboard does.
- **Typography.** A three-typeface system: IBM Plex Sans Condensed for headings, the wordmark, and section labels; IBM Plex Mono — with tabular numerals — for every raw technical value (hashes, IPs, timestamps, scores); body text left untouched to avoid any risk to the existing dense data tables.
- **Iconography and mark.** An original, wholly geometric logo (a horizon line with two open "signal" nodes converging on one filled "detection" node) that literally depicts the tagline rather than just looking generically tactical, with no insignia, seal, shield, or existing company/military symbol of any kind.

Where It Landed

The identity was applied first to the shared components that appear everywhere (the header, cards, buttons, inputs, badges) so it took effect consistently across every page from one set of edits, then rolled out page-by-page to the ten pages sharing that header/nav, and finally into every exported report format (PDF, CSV, Markdown), each of which now carries the platform name, a generation timestamp, and an Investigation ID. The result is documented with real, live screenshots throughout this submission rather than mockups.




Figure 6 — The login screen — the mark, wordmark, tagline, and the subtle grid-and-horizon-line background introduced in the branding pass.

Brand Identity

Name and Tagline

HORIZON GRID — "Every Signal. One Operational Picture."

The name and tagline are meant to be read literally, not just atmospherically: "signal" is what a single provider or evidence item contributes; "grid" and "horizon" evoke a fixed reference frame that many independent signals get plotted onto; "one operational picture" is the actual product outcome — a correlated, scored, explainable view assembled from everything the platform collected, rather than a raw dump of provider responses.

Logo Concept

The mark is wholly geometric and original: a horizon line with two open "signal" nodes converging on one filled "detection" node. It is a literal depiction of the tagline — multiple signals resolving into one picture — rather than a generic shield, seal, or insignia shape. It exists in three forms:

- **Full lockup** — the mark plus the "HORIZON GRID" wordmark, used on the login/register screens and the About page.
- **Compact lockup** — the mark plus "HG" set inside a corner-bracket badge frame, used in the persistent header on every authenticated page.
- **Favicon** — a three-primitive degraded form of the same mark, simplified to stay legible at 16px.

No part of the mark reuses or resembles any existing military insignia, government seal, NATO symbol, or real defense-contractor logo (Lockheed Martin or otherwise) — a constraint that was explicit in this pass's own brief and was checked directly against the final design.

Color System

A CRT-phosphor teal-cyan is the primary color, chosen specifically to avoid the generic SaaS-dashboard blue most competing tools default to. It sits against a near-black "anodized steel panel" background with a warm off-white foreground — dark enough to read as an operational console, warm enough to avoid the colder, more generic near-black-and-electric-blue palette associated with stock "hacker" UI kits.

Status Indicator Language

Every operational status in the product — provider health, system health, investigation state — is expressed through one shared six-state vocabulary: **OPERATIONAL, DEGRADED, WARNING, CRITICAL, OFFLINE, UNKNOWN**. Each state is encoded through a distinct hue, fill density, border style, *and* icon simultaneously, so the state is never carried by color alone. This same vocabulary is what Provider Health's four real status categories (Healthy/Degraded/Down/Unknown) map onto — mapped honestly (Healthy→Operational, Down→Offline) rather than inventing new states for data the backend doesn't actually report.

Typography

- **IBM Plex Sans Condensed** — headings, the wordmark, section/panel labels.
- **IBM Plex Mono**, with tabular numerals — every raw technical value: IP addresses, hashes, timestamps, coordinates, counts, and scores.
- **Inter** (unchanged) — body copy, left untouched to avoid disturbing the product's existing dense data tables.

Shape and Elevation

Outer containers (panels, cards) keep a softer corner radius; interactive and status elements (buttons, inputs, badges, chips) use a visibly sharper radius — a deliberate contrast meant to read as soft panels precisely containing exact controls. Cards are a flat fill bounded by a hairline border rather than drop-shadowed, so the interface reads as machined layers rather than a floating card stack.

Navigation and Panels

The persistent header (`BrandHeader`) pairs the compact logo lockup with a live system-status indicator sourced from the real (`GET /health/detailed`) endpoint (polled every 60 seconds, never hardcoded), the global search bar, and the existing workspace navigation grouped by function (Command / Intelligence / Analysis / Operations / Administration). Section headers throughout the product (`CardTitle`) render as uppercase, letter-spaced, muted "instrument-panel" labels by default — an "operational readout" register applied once at the shared-component level so it took effect consistently everywhere without a per-page edit.




Figure 7 — The About page — the full mark lookup, live version and dependency status, and the platform/documentation summary.




Figure 8 — The compact mark and status language reused in context — the AI backend selector on the home page, showing all eleven configured backends.

Why This Design Was Chosen

The identity had to satisfy two requirements that are easy to satisfy separately but hard to satisfy together: look like professional, mission-ready operational software, and be demonstrably original — no real insignia, seal, or defense-contractor mark anywhere in it, and no cyberpunk-movie neon-glow excess either. "Datum Signal" was the strongest of three independently generated candidate identities, chosen because its status language and shape language carry the "operational instrument" feeling structurally (through contrast, iconography, and a consistent status vocabulary) rather than through any single borrowed visual trope.

Bug History: What Actually Broke, and What Got Fixed

This project's own QA and hardening reports are written as adversarial, live-system audits rather than marketing copy, and several of them explicitly disqualify the product from a "release ready" verdict at the time they were written. That candor is treated here as material worth keeping, not smoothing over. This section walks through five phases of the project's development — early hardening, the AI-backend expansion, the Security Assessment Toolkit / port scanner, the mission-critical reliability pass, and the final release cycle — and lists the real bugs each phase's own reports document, including several that were found and disclosed but never fixed. Every row below is sourced from a specific QA/report document or a specific source file read directly; where a claim could not be independently re-verified against code, that is stated rather than implied.

Early Hardening

The earliest hardening work covered in the project's reports is not glamorous: two divergent, undocumented backend virtualenvs, a broken credential-test flow in the Windows Setup Wizard, a `.env` writer with no escaping, and a first pass at multi-admin RBAC and database-backed runtime provider configuration. Threaded through this phase is `ENTERPRISE_QA_PHASE1_FINDINGS.md`, a large adversarial red-team pass run directly against the live, running stack — it is explicitly an audit report, not a fix report, and states plainly that it found problems without fixing them. Only a representative subset of its 74 findings is reproduced below; the rest remain in that document, disclosed and open at the time it was written.

Bug	Impact	Root Cause	Fix	Status
Two divergent, drifted backend virtualenvs (<code>backend/.venv</code> , <code>backend/.venv_test</code>)	<code>pytest app/tests/unit</code> gave 4 collection errors, 0 tests run (<code>ModuleNotFoundError: No module named 'cryptography'</code>)	No documented canonical venv; <code>asynpg==0.29.0</code> had no Windows/Python-3.13 wheel, forcing ad hoc workarounds per venv	Deleted <code>backend/.venv</code> ; bumped <code>asynpg</code> to <code>0.30.0</code> ; reinstalled <code>.venv_test</code> from <code>requirements.txt</code> ; declared it canonical in <code>docs/TESTING.md</code>	Fixed and verified — 282/282 unit tests passing, 5 consecutive clean runs
Fresh <code>docker compose up -d postgres redis</code> leaves an empty DB schema	Running the suite without a manual <code>alembic upgrade head</code> produced 36 failed + 27 errors, all "relation does not exist" — looked like an app bug but was a missing setup step	Migrations only auto-run when the <code>backend</code> service itself starts	Documented the exact <code>alembic upgrade head</code> command in <code>docs/TESTING.md</code>	Fixed (documentation only)
Setup Wizard "Test Connection" required a session token only ever set inside the "Start Installation" handler	On a reconfigure run, no login ever happened, so every Test Connection click failed with a misleading "create the administrator account first" error even though the account already existed	Token-acquisition logic lived only in the install-click handler, not as an independent path	New <code>GetOrCreateWizardSession</code> (login-only, never registers an account), called directly from each Test Connection click	Fixed
<code>Write-EnvFile.ps1</code> wrote <code>.env</code> values with no escaping	A <code>#</code> in a password/API key silently truncated the rest of the line; embedded newlines could corrupt the file	No escaping/validation before writing values to <code>.env</code>	New <code>ConvertTo-SafeEnvValue</code> strips CR/LF and rejects embedded <code>#</code>	Fixed
Hybrid Analysis provider returned HTTP 301 instead of a valid result	Real Hybrid Analysis lookups were silently broken, not just the test button	<code>httpx</code> doesn't follow redirects by default; <code>hybrid-analysis.com</code> now permanently redirects to the bare domain	Fixed in <code>app/providers/connection_test.py</code> and <code>app/providers/stubs/hybrid_analysis.py</code> , plus the unit-test mock URL	Fixed and confirmed live (243ms, correct <code>no_data</code> for a test hash)
Admin-initiated password reset left every already-issued access/refresh token valid	The opposite of the intended security effect — a compromised session stayed alive after "resetting" the password, for up to 7 days	JWTs are stateless (signature + expiry only); nothing tied token validity to password state	Added <code>token_version</code> to <code>User</code> ; <code>get_current_user()</code> and <code>POST /auth/refresh</code> reject any token whose embedded version doesn't match the DB value; <code>reset_password()</code> increments it	Fixed, verified live and by automated test
<code>/auth/me</code> always returned <code>full_name:</code> ""	Display bug regardless of stored value	Found and fixed as part of the RBAC work	Corrected to return the stored value	Fixed
2 pre-existing tests broke against the new async, tuple-returning <code>_get_ai_client()</code> signature	Regression in the runtime-provider work	Signature change not reflected in test doubles	Caught and fixed in the same session's regression pass	Fixed — 153/153 passing afterward
Oversized IOC value (e.g. a long URL) crashes the <code>/lookup/stream</code> SSE generator, leaving the investigation stuck at <code>status='running'</code> forever	Any user with ordinary <code>lookup:create</code> permission can corrupt a lookup into an unrecoverable state by pasting a long value	OTX's <code>source_url</code> has no length cap and can exceed the <code>provider_results.source_url VARCHAR(2048)</code> column; the generator's exception handler then does a second <code>commit()</code> on an already-rolled-back session, raising an uncaught <code>PendingRollbackError</code> that escapes the "mark FAILED" safety net entirely	None found in this report	Open (disclosed, not fixed)
No request-correlation ID anywhere in the backend	Concurrent operations are unattributable from <code>docker logs</code> alone	<code>app/main.py</code> registers only <code>CORSMiddleware</code> and the Prometheus instrumentator; no request-ID middleware anywhere	None found in this report	Open (disclosed, not fixed)
Failed logins are invisible in stdout logs beyond a bare <code>401</code>	Security-relevant events not surfaced in ordinary log inspection	<code>login()</code> in <code>app/api/routes/auth.py</code> has zero <code>logger.*</code> calls	None found in this report	Open (disclosed, not fixed)
Login response time leaks account existence via a ~26x timing side-channel	Existing-email+wrong-password median 231.4ms vs. nonexistent-email median 8.8ms (N=40/arm)	<code>if not user or not verify_password(...)</code> short-circuits — <code>bcrypt</code> only runs for real emails	None found in this report	Open (disclosed, not fixed)
No rate limiting/backoff/lockout on <code>/auth/login</code>	15 rapid failed logins all returned 401 with no throttling signal	<code>lookup.py</code> has an explicit <code>Ratelimiter</code> ; <code>/auth/login</code> does not	None found in this report	Open (disclosed, not fixed)
Weak, hardcoded default DB credentials baked into <code>backend/app/core/config.py</code> // <code>docker-compose.yml</code>	Any deployment that skips the wizard's random-secret generation and runs bare <code>docker compose up</code> is exposed	Default <code>postgresql+asynpg://ioc:ioc@postgres:5432/ioc_intel</code> ; live-confirmed the running Postgres instance actually used this default	None found in this report; mitigated only by the wizard, not the app itself	Open (disclosed, not fixed)
<code>jwt_secret_key</code> has a hardcoded fallback ("change-me-in-production") with no startup guard	Same string signs every JWT and, via HKDF, derives the at-rest credential-encryption key	No code-level safety net for a deployment path that skips the wizard	None found in this report; confirmed not the live value in practice because the wizard generates a real secret	Open (disclosed, not fixed)

The QA phase-1 report also discloses a large number of P2/P3/P4 findings not reproduced in the table above — among them a stale AI-provider "last test: failed" badge that can persist on a healthy backend, `mitre_mappings` coming back empty for a textbook Log4Shell case, an unhandled `IntegrityError` on concurrent duplicate "Add to Basket" requests, CSV/PDF export-injection issues (carried forward and addressed later — see the Final Release phase below), and a `pip-audit` false negative against packages independently confirmed vulnerable via direct OSV.dev queries. All of these are recorded in `ENTERPRISE_QA_PHASE1_FINDINGS.md` as open at the time of writing; this document does not claim any of them were fixed within that report.

AI Backend Expansion

The provider registry for AI backends (`backend/app/core/runtime_config.py` line 29) lists 11 entries: `ollama`, `anthropic`, `bedrock`, `gemini`, `groq`, `openai`, `kimi`, `deepseek`, `xai`, `mistral`, `openrouter`. Per `HORIZON_GRID_AI_PROVIDER_EXPANSION_QA_REPORT.md`, six of those (Ollama, Anthropic, Bedrock, Gemini, Groq, OpenAI) predate this expansion; five (Kimi, DeepSeek, xAI/Grok, Mistral, OpenRouter) were added in this v0.2.0 pass, each reportedly given a live base-URL check, forced `tool_choice` verification, live model discovery, and a dedicated batch of new unit tests. The regression suite is reported at 313 passed (282 pre-existing + 31 new). None of the live HTTP timing figures, the 313-test total, or the installer hash claims in that report were independently reproduced this turn — they are relayed as the report's own claims, not independently confirmed.

The report's own "fixed defects" section lists exactly three bugs found during this pass, plus one carried-forward item explicitly flagged as not new to this cycle:

Bug	Impact	Root Cause	Fix	Status
<code>python-jose 3.4.0 + pyasn1 0.6.4</code> <code>pip ResolutionImpossible</code>	Dependency install failure	Version conflict between the two pinned packages	Resolved in the immediately preceding backend-test-stability cycle; re-verified still holding in this release	Fixed (carried forward, not new to this pass)
Windows installer bundled the local dev virtualenv	The installer bloated from ~69 MB to bundling ~190 MB of unused dev-venv content	<code>installer.iss</code> never excluded <code>backend/.venv / .venv_test</code>	Excluded the dev venv directories from the installer build	Fixed
<code>standalone-ai-guide.md</code> never documented the OpenAI backend	OpenAI, added as the 6th backend in an earlier cycle, had no user-facing documentation until this pass	Documentation not updated when OpenAI was originally added	Documentation updated during this pass	Fixed

Two further items are recorded as corrections rather than bugs: an earlier, narrower fetch of a Mistral guide page had incorrectly suggested named-tool forcing wasn't supported, corrected against the authoritative API reference; and a documented (not fixed, because it is not a defect) Moonshot/Kimi API behavior — forcing `tool_choice` conflicts with Kimi's "thinking" models (`kimi-k3`, `kimi-k2.7-code`), returning HTTP 400 — was worked around by choosing `kimi-k2.5` as the default model, with a dedicated test (`test_kimi_400_thinking_conflict_is_reported_clearly`) asserting the error surfaces clearly rather than silently. `backend/app/ai/service.py` contains inline comments describing several other real, previously-fixed bugs (a Spamhaus `"listed": false` misread, an NVD-field-length prompt-pollution issue, a self-contradictory verdict/probability bug, an EICAR-hash fabrication bug), but nothing in the expansion QA report ties those specifically to the 6→11 backend expansion, so they are not attributed to this phase here.

Security Assessment Toolkit (Port Scanner)

The Security Assessment Toolkit wraps `nmap` behind three fixed scan profiles (`quick`, `standard`, `web`) with no caller-constructed argv strings — targets are passed as a single, always-terminal `asyncio.create_subprocess_exec` argument, never through a shell (`backend/app/security_assessment/nmap_tool.py`). Two dedicated QA reports, `HORIZON_GRID_PORT_SCANNING_QA_REPORT.md` and `SECURITY_ASSESSMENT_QA_REPORT.md`, document a cluster of real bugs found while building cancellation support and hardening this feature across versions v0.2.1 and v0.2.2. A subset of these — the cancellation flow itself, the IPv6 `-6` flag fix, the profile-id validation, the malformed-CIDR handling, and the completion-status race guard — were independently confirmed this turn by reading `backend/app/security_assessment/nmap_tool.py` and `backend/app/core/security_assessment.py` directly; the remainder are relayed as the QA reports' own claims, marked accordingly.

Bug	Impact	Root Cause	Fix	Status
No way to cancel a running scan	A slow or regretted scan could only be waited out; restarting the backend to escape it would orphan the running <code>nmap</code> process	No cancel endpoint, button, or subprocess-kill path existed at all	Added <code>CANCELLED</code> status, a cancel endpoint, a task-tracking dict, dedicated <code>CancelledError</code> handling in both the orchestrator and the nmap tool, and restart-resilience cleanup	Fixed (v0.2.1) — confirmed present in <code>security_assessment.py</code> (<code>cancel_run()</code> , <code>_execute_run()</code>) and <code>nmap_tool.py</code> 's <code>except asyncio.CancelledError</code> branch
IPv6 scans never actually probed the target	High impact — IPv6 scanning was silently non-functional; nmap exited 0 having scanned 0 hosts, indistinguishable from a real clean scan	<code>nmap</code> needs an explicit <code>-6</code> flag for an IPv6 literal; it was missing	Added <code>-6</code> conditionally for IPv6 targets	Fixed (v0.2.2, found by an independent re-verification pass) — confirmed present at <code>nmap_tool.py</code> lines 124–131
Migration added the enum value <code>'cancelled'</code> lowercase against Postgres's existing uppercase enum labels	Two tests failed with <code>invalid input value for enum securityassessmentrunstatus: "CANCELLED"</code>	Case mismatch between the new migration and existing enum labels	Changed to <code>ADD VALUE IF NOT EXISTS 'CANCELLED'</code>	Fixed
Two concurrent cancel requests for the same run could each write their own audit row	Audit trail showed a single cancellation as two events (audit-accuracy issue, not a functional one)	No single designated writer for the audit event on the live-task cancellation path	Made the run's own background task the sole writer for the live-task path; the restart-orphan path gates its write on rowcount	Fixed — confirmed present at <code>security_assessment.py</code> 's <code>cancel_run()</code> and <code>_execute_run()</code>
An unknown or mistyped scan profile id was silently accepted and reported as a successful, empty scan	High impact — indistinguishable from a real clean scan, with <code>error_message: null</code>	Profile id not validated before creating the run row	Validate the profile id against the tool's known profiles up front, returning 400	Fixed — confirmed present at <code>security_assessment.py</code> lines 298–313
A malformed CIDR lookup value crashed the run endpoint with a raw 500	Medium impact — unhandled server error on bad input	<code>ipaddress</code> parsing raised an uncaught <code>ValueError</code>	Catch the <code>ValueError</code> , return 400	Fixed — confirmed present at <code>security_assessment.py</code> lines 260–266
A run's own successful completion could clobber a status another writer had already finalized	Medium impact — a run could end up <code>COMPLETED</code> while still carrying a stale <code>FAILED</code> error message	No guard against a race between the run's own completion write and a concurrent finalizer (e.g. startup orphan-recovery)	Guard the completion status update with <code>WHERE status IN (PENDING, RUNNING)</code>	Fixed — confirmed present at <code>security_assessment.py</code> lines 462–474
Empty <code>tool_ids</code> was accepted, producing a no-op run	Low impact	No validation on an empty tool selection	Reject at the schema level (422)	Fixed per the report; not independently verified in the schema file this turn
Spamhaus provider misreported a DNS query-rejection as a positive "malicious" verdict	A benign domain (e.g. <code>example.org</code>) could be flagged malicious for no real reason	Query-rejection response mishandled as a positive listing	Correctly report <code>unknown</code> on rejection	Fixed per the report (found incidentally during regression, unrelated to scanning); not independently verified this turn
Confusing generic 401 instead of 403 "Account disabled" on the losing side of a concurrent last-admin-protection race	Misleading error on an edge-case concurrency path	Race condition in last-admin-protection logic	Corrected error response	Fixed per the report; not independently verified this turn
An AI-generated assessment could cite findings in prose while leaving the structured evidence list empty	UI clickable-citation feature broken for that case	No schema validator enforcing citation/evidence-list agreement	New schema validator using the existing retry-on-validation-failure mechanism	Fixed per the report; not independently verified this turn
Two new regression tests failed in the CI "unit" job (no <code>nmap</code> binary there)	CI-only test failure	Mock didn't cover <code>is_available()</code>	Mock <code>is_available()</code> too	Fixed per the report
Test-isolation race: <code>SecurityAssessmentRun.status</code> flips to <code>COMPLETED</code> before the post-completion AI-refresh step finishes	Flaky/corrupting test suite — a test polling only the status field could tear down its shared DB engine mid-refresh	No real synchronization primitive between test assertions and the in-flight background task	Added <code>wait_for_background_runs()</code> , awaiting the actual in-flight tasks	Fixed — confirmed present at <code>security_assessment.py</code> lines 168–177

Two disclosed, non-bug design gaps from this phase are worth carrying forward for completeness: `POST /run` accepts one shared `profile` string for every selected tool, but only `nmap` defines more than one profile, so selecting multiple tools with an nmap-specific profile like `"quick"` would make every non-nmap tool reject it — handled in the frontend by only offering profile ids valid across the intersection of every selected tool's profile set, per the report (not independently verified in the frontend code this turn). Separately, `SECURITY_ASSESSMENT_QA_REPORT.md` and `HORIZON_GRID_PORT_SCANNING_QA_REPORT.md` both disclose that no live elevated Windows installer run or live Linux install/systemd run was performed for either version tested, and that TLS-downgrade testing and DNS subdomain enumeration remain out of scope by design.

Mission-Critical Hardening

`MISSION_CRITICAL_CERTIFICATION_REPORT.md` certifies v0.2.3 specifically, and states the version number explicitly at line 4. Its stated purpose was unattended-remote-deployment reliability: installer gating, restart/reboot recovery, automatic service recovery, health checks, backup/restore, and a 3-hour, 171-cycle soak test against the live stack (memory reported flat throughout, zero spontaneous container restarts). Two other files in this same research bundle, `LAN_DEPLOYMENT_QA_REPORT.md` and `LAN_REGRESSION_TEST_REPORT.md`, are dated earlier (2026-08-12/14) and cover a separate, distinct effort — moving the platform from localhost-only to LAN accessibility — and carry no version number of their own; one bug below is sourced from that separate work and is labeled as such rather than folded into v0.2.3.

Bug	Impact	Root Cause	Fix	Status
Linux <code>horizon-grid.service</code> systemd unit was defined but never actually <code>systemctl enable</code> d anywhere in the codebase	A working install would silently not survive a host reboot	The <code>[Install]</code> section declared <code>WantedBy=multi-user.target</code> , but only <code>daemon-reload</code> ran in <code>postinst</code> — <code>enable</code> was never called	Added the enable call to the wizard's success path	Fixed and verified live via <code>systemctl is-enabled</code> → <code>enabled</code>
<code>check-prerequisites.sh</code> 's JSON output had a permanently empty <code>checks</code> array	Structured output from the prerequisite gate was silently useless (console output and exit code were unaffected, computed separately, which is why it went unnoticed)	Bash's lowercase <code>true</code> / <code>false</code> were interpolated as bare Python identifiers, causing a silent <code>NameError</code> on every call, swallowed by an error fallback	Passed fields through the environment instead of interpolating into generated source	Fixed; verified live for both pass and simulated-failure cases
<code>BaseProvider.run()</code> swallowed exception types needed by the orchestrator's retry loop	<code>provider_max_retries</code> had no effect for providers without their own network-error handling	Exception handling too broad/incorrect at the base-provider layer	Corrected exception handling to let retryable errors propagate to the orchestrator's retry loop	Fixed, live-verified with a fake provider showing a <code>ConnectError</code> now genuinely retried
Final Assessment panel showed an identical badge for a genuine AI failure and a correct no-evidence skip	Users could not distinguish "AI backend is down" from "there was nothing to assess"	Backend already tracked/sent <code>ai_outcome</code> ; the frontend never read it	Frontend updated to read and render <code>ai_outcome</code>	Fixed, confirmed against a real <code>ai_outcome="failed"</code> record
<code>docker-compose.yml</code> 's <code>PUBLIC_API_URL</code> fallback used <code> \${PUBLIC_API_URL:-http://localhost:8000}</code> (colon-hyphen)	Fell back to the hardcoded default even when the variable was explicitly set to empty	Colon-hyphen bash/compose substitution treats an explicitly-empty value the same as unset	Changed to <code> \${PUBLIC_API_URL-}</code>	Fixed — from the separate LAN-deployment work (<code>LAN_DEPLOYMENT_QA_REPORT.md</code>), not part of v0.2.3 itself
A full-suite regression run reported 365 passed, matching the previous known-good count rather than reflecting the session's actual changes	Test methodology error — the suite had validated a stale artifact, not the current code	The <code>hgmc</code> test container was running a stale Docker image predating the session's backend changes	Rebuilt the image, re-ran the suite for real	Fixed — 380 passed after rebuild
Redis did not persist across restarts	Cache/session state lost on every restart	Not detailed beyond the fact itself in the report	Redis now persists across restarts	Fixed, confirmed live

The certification report is explicit about what it did not fix. Disclosed and left open at the time of writing: no automated host-disk-space alerting; no retention/cleanup job for growing investigation tables; Neo4j and OpenSearch fully provisioned with zero actual read/write traffic (an open resource-cost question, not a bug); a DNS-rebinding TOCTOU gap on the SSRF check (it validates a DNS snapshot, but the real outbound call re-resolves independently); no off-site/off-host backup copy; no global ONLINE/LIMITED/OFFLINE UI indicator; no real simulated power-loss test performed; per-subprocess (nmap) OS-level resource limits remain container-level only; `structlog` is configured but unused anywhere in application code; the full ~40-document existing doc set was not exhaustively regenerated; and a live elevated Windows installer run was not performed in this session — the report gives the specific, disclosed reason that it required an interactive UAC prompt the autonomous session could not satisfy. The report's own stated verdict is "**MISSION-CRITICAL READY WITH DOCUMENTED LIMITATIONS**," explicitly not an unqualified "ready" because of that limitations list, and explicitly not "not ready" because every one of the gaps it says it found in its own review was fixed and verified.

Final Release

The final-release phase in this bundle spans at least two distinct product states and dates, and the six source files should not be read as one continuous story. `FINAL_QA_REPORT.md` (2026-08-14, IOC Intelligence Platform) is a red-team regression pass that states a top-line verdict of **NOT RELEASE READY**, citing a credential-wipe bug, an apparently non-functional export feature, and plaintext secrets exposure as each independently disqualifying per that mission's own severity rules. `RELEASE_BLOCKER_TRACKER.md` tracks the same bugs through fix-test-break-test-again cycles for the same day. `FINAL_SHIP_AUDIT_REPORT.md` , also 2026-08-14, is a separate release-candidate lock/audit (not a fresh QA pass) that reaches **RELEASE READY WITH KNOWN LIMITATIONS**, with a same-day addendum that closed two of its own disclosed limitations outright. `FINAL_RELEASE_QA_REPORT.md` documents a later, separate effort — the "HORIZON GRID" rebrand — also concluding **RELEASE READY WITH KNOWN LIMITATIONS**. The most recent file by timestamp, `FULL_FUNCTIONAL_TEST_REPORT.md` (2026-08-21, v0.2.3), is explicitly marked "Status: IN PROGRESS" and states its own verdict is interim, not final, covering roughly a third of a planned 33-phase mission. No file in this bundle reconciles these into one single authoritative final verdict for the product as a whole.

Bug	Impact	Root Cause	Fix	Status
Credential-wipe-on-save (BUG-01)	Saving certain configuration silently wiped stored credentials	Not detailed beyond the bug identifier in this bundle's summary of the report	First round of fixes was found by adversarial re-check to have gaps; a second round added row-locking and real DB-backed regression tests	Fixed (<code>FINAL_QA_REPORT.md</code> → <code>RELEASE_BLOCKER_TRACKER.md</code>), required two rounds
Silent AI-assessment failure on a self-contradictory verdict (BUG-02)	An AI-generated assessment could fail silently rather than surfacing an error	Not detailed beyond the bug identifier in this bundle's summary of the report	Fixed in the same cycle	Fixed (<code>RELEASE_BLOCKER_TRACKER.md</code>)
Unbounded provider data burying severity signal in the AI prompt (BUG-03)	Severity signal could be lost in an oversized prompt	Not detailed beyond the bug identifier in this bundle's summary of the report	Fix caused two further regressions in its own correction rounds: first broke MITRE ATT&CK's long descriptions, then broke Groq's request-size limit; both were subsequently corrected	Fixed, after two additional correction rounds
Export feature reported as apparently completely non-functional (BUG-04)	Would have been release-blocking if true	Original test methodology was wrong — the JSON/Markdown export path is client-side and never hit the backend route being tested	N/A — retracted	Retracted, not a real bug — only PDF/CSV correctly 404'd, which was already-documented, non-regression behavior at that time
PDF/CSV export not implemented	Two export formats unavailable at initial ship-audit time	Feature not yet built	Implemented in a same-day addendum to <code>FINAL_SHIP_AUDIT_REPORT.md</code> ; new installer build produced (180 tests passing)	Fixed (same-day addendum)
A narrow first-time-configure concurrency edge case in the credential row-lock fix	Edge-case race during first-time configuration	Related to the row-locking fix applied for BUG-01	Closed in the same same-day addendum	Fixed (same-day addendum)
Windows installer password-mismatch crash	Real installer failure during the HORIZON GRID rebrand cycle	Not detailed beyond the fact of the fix in this bundle's summary	Fixed as part of <code>FINAL_RELEASE_QA_REPORT.md</code> 's work, with a user-confirmed real installer run	Fixed
Stale shipped <code>.exe</code> missing <code>docker-compose.yml</code> (P0)	Installer would fail to bring up the stack correctly	Packaging step did not include the compose file in the shipped installer	Not detailed further in this bundle beyond the fact that it was found and fixed	Fixed (<code>FULL_FUNCTIONAL_TEST_REPORT.md</code>)
Leftover registry <code>InstallLocation</code> from an old test harness silently redirecting installs to a fake path (P0)	Installs could be silently redirected to an incorrect location	Stale registry key left behind by a prior test harness	Not detailed further in this bundle beyond the fact that it was found and fixed	Fixed (<code>FULL_FUNCTIONAL_TEST_REPORT.md</code>)
Free-text <code>threat_assessment</code> field able to contradict the validated <code>final_verdict</code> (P1)	An AI-generated assessment's narrative text could disagree with its own structured verdict	No validator enforcing agreement between the free-text field and the structured verdict at the time this was found	Not detailed further in this bundle beyond the fact that it was found and fixed	Fixed (<code>FULL_FUNCTIONAL_TEST_REPORT.md</code>)

Disclosed limitations from this phase, not treated as blockers by the reports that state them: Groq's free/shared-tier rate limit (12,000 tokens/minute) is a real external constraint, handled cleanly rather than retried into a storm; AI-generated verdict quality is model-dependent, most visibly with smaller local models; a live plaintext-secrets exposure (including an AWS key pair and several threat-intel provider keys) was found in a dev-environment `.env` file and flagged for rotation, but was not itself rotated by the QA process; and a disposable QA test account and test case were flagged for removal before handing the environment to a real end user. `FULL_FUNCTIONAL_TEST_REPORT.md` additionally lists nine planned phases — log rotation, performance under load, a broader security/RBAC sweep, a Windows reboot test (deliberately deferred), a Linux install-and-reboot test, uninstall/reinstall, full regression, a soak test, and a final post-fix repeat of the critical path — as explicitly not yet run rather than silently assumed to pass, consistent with its own stated interim, not final, verdict.

The Testing Journey

Scope and Sourcing Note

This chapter reconstructs the QA history of the IOC Intelligence Platform / HORIZON GRID codebase from the project's own dated QA and hardening reports, plus direct reads of the source files those reports describe (`backend/app/core/runtime_config.py` , `backend/app/security_assessment/nmap_tool.py` , `backend/app/core/security_assessment.py` , `backend/app/scoring/engine.py` , `backend/app/ai/service.py` , `backend/app/ai/schemas.py` , `backend/app/models/user.py` , `backend/app/auth/rbac.py` , `backend/app/api/routes/admin.py` , `backend/app/providers/registry.py` , `backend/app/core/dashboard.py` , and related route/frontend files). Where a claim comes only from a report's own text rather than from a source file that was independently opened and read, that distinction is preserved below rather than smoothed over. Several reports disclose their own unverified or unreproduced claims; those disclosures are kept intact rather than upgraded into confirmed fact. The project spans at least two version lines and a rebrand — "IOC Intelligence Platform" v0.1.0 and its later identity as "HORIZON GRID," continuing through v0.2.0, v0.2.1, v0.2.2, and v0.2.3 — and the reports below are attributed to the specific version and date each one actually covers, never blended into a single number or a single verdict that no source document itself states.

Early Hardening: Environment Drift and Configuration Bugs

The earliest layer of QA work documented is not feature testing but environment and tooling hardening. `BACKEND_TEST_ROOT_CAUSE_REPORT.md` found that the backend had accumulated two divergent, undocumented virtualenvs (`backend/.venv` and `backend/.venv_test`), one of which had drifted badly from `requirements.txt` — `cryptography` was missing entirely, and `fastapi` / `pytest` were both several versions ahead of their pins. Running `pytest app/tests/unit` against the drifted environment produced 4 collection errors and 0 tests run. The root cause traced back further: `asynccpg==0.29.0` , the original pin, has no Windows/Python-3.13 wheel, which had forced whoever set up each venv to improvise a workaround, or skip it, independently. The fix deleted the stale `.venv` , bumped `asynccpg` to `0.30.0` (the first version with a `cp313` Windows wheel), reinstalled `.venv_test` cleanly from `requirements.txt` , and declared `.venv_test` canonical in `docs/TESTING.md` . Verified result: 282/282 unit tests passing, 0/32 pin mismatches, across 5 consecutive clean runs on both native Windows and a `python:3.12-slim` container. The same report separately found that `docker compose up -d postgres redis` alone leaves an empty database schema — migrations only run automatically when the `backend` service itself starts — which had produced 36 failed and 27 errors that looked like application bugs but were actually a missing `alembic upgrade head` step; this was fixed by documenting the step, not by changing code. One pre-existing issue was disclosed and explicitly left unfixed in this pass: an "Event loop is closed" cross-file integration-test isolation problem (SQLAlchemy/Redis pools bound to a since-closed per-test event loop), producing 4 failures and 6 errors when the full integration suite runs outside the container-based CI path — noted as not affecting the unit suite that actually gates CI.

`API_CONFIGURATION_FIX_REPORT.md` covers a genuinely broken part of the Windows Setup Wizard: the "Test Connection" buttons required a session token that was only ever set inside the Summary page's "Start Installation" handler, so on a reconfigure run (a blank Admin Account page, by design) no login ever happened and every Test Connection click failed with a misleading "create the administrator account first" error even though the account and platform already existed. The fix added a dedicated, login-only `GetOrCreateWizardSession` function callable directly from each Test Connection click, independent of the install path. The same report found that `Write-EnvFile.ps1` wrote `.env` values with no escaping at all — a `#` in a password or API key silently truncated the rest of the line (read as a comment), and embedded newlines could corrupt the file outright, a real, previously invisible cause of "the key I entered doesn't work after saving." The fix added `ConvertTo-SafeEnvValue` , which strips CR/LF and rejects embedded `#` before any value is written. It also found that the Hybrid Analysis provider was returning HTTP 301 instead of a real result, because `httpx` does not follow redirects by default and `hybrid-analysis.com` 's API now permanently redirects to a bare domain — fixed in both `connection_test.py` and the provider stub, and confirmed live (243ms, correct `no_data` for a test hash). A cosmetic UI overlap in the Groq panel of the Setup Wizard (the API-key field overlapping the model dropdown) was also fixed. Separately, the report discloses a real secret leak that was not a code bug: a scratch UI-automation log outside the repo had captured a live Groq API key in plaintext, redacted immediately on discovery — and flags, without acting on it, that no git repository existed at the time, so `.gitignore` provided no actual enforcement against a plain folder-copy leaking a real `.env` .

`ADMIN_RBAC_QA_REPORT.md` documents both the build-out of a genuine multi-admin RBAC console and a real security bug caught along the way: an admin-initiated password reset changed a user's `hashed_password` but left every already-issued access and refresh token valid until natural expiry (up to 7 days for refresh tokens) — the opposite of what a security-motivated password reset should do. The fix added a `token_version` column to `User` , incremented on every reset; `get_current_user()` and `POST /auth/refresh` both now reject any token whose embedded `token_version` doesn't match the current database value (confirmed directly in `backend/app/auth/rbac.py` , lines 63-69, and in `backend/app/models/user.py` 's column comment). This is verified both live and by an automated regression test. The same report incidentally found, and left as pre-existing and out of scope, a home-page `AiQuickSwitch` widget calling an admin-only endpoint regardless of the signed-in user's role (producing a benign 403 in the browser console for non-admins), and fixed, as part of this same pass, a bug where `/auth/me` always returned an empty `full_name` regardless of the stored value.

`RUNTIME_PROVIDER_IMPLEMENTATION_REPORT.md` covers the move from a frozen `.env` -at-startup credential model to a database-backed, Fernet-encrypted, runtime-mutable one. It discloses a design tradeoff rather than a bug: `ENCRYPTION_MASTER_KEY` defaults to an HKDF derivation from `jwt_secret_key` , so rotating the JWT secret would silently break decryption of previously stored provider credentials unless a separate master key is explicitly set — documented, not enforced. During this same work, 2 pre-existing tests broke against a new async, tuple-returning `_get_ai_client()` signature; both were caught and fixed within the same session, leaving 153/153 tests passing afterward.

Regression Testing: Expanding the AI and Provider Surface

Two later missions specifically exercised regression coverage as the platform's provider and AI surface grew, and the platform's `backend/app/scoring/` engine added a third layer of coverage focused purely on scoring correctness.

`HORIZON_GRID_AI_PROVIDER_EXPANSION_QA_REPORT.md` documents the v0.2.0 expansion of AI backends from 6 to 11. Confirmed directly against `backend/app/core/runtime_config.py` line 29, the live `AI_BACKENDS` list is `["ollama", "anthropic", "bedrock", "gemini", "groq", "openai", "kimi", "deepseek", "xai", "mistral", "openrouter"]` — 11 entries, consistent with `_AI_ENV_SEED_MAP` (lines 68-88) and `backend/app/ai/service.py` 's `_model_id_for_backend()` / `_build_client()` . Per the report itself (not independently reproduced by this reading): Ollama, Anthropic, Bedrock, Gemini, Groq, and OpenAI predate this pass; Kimi, DeepSeek, xAI, Mistral, and OpenRouter were newly added, each with a live base-URL check, forced `tool_choice` verification, live model discovery, and its own unit tests (Kimi 7, DeepSeek 6, xAI 6, Mistral 6, OpenRouter 6). The report states a regression total of 313 passed (282 pre-existing plus 31 new), with "Failed tests: None." Three real defects are named as found and fixed during this pass: a carried-forward (not newly found) `python-jose` / `pyasn1` pip dependency-resolution conflict, re-verified as still holding; the Windows installer bundling the entire local dev virtualenv (`installer.iss` never excluded `backend/.venv` / `.venv_test` , inflating the installer by roughly 190 MB of unused content), fixed during this release's installer rebuild; and a documentation gap where the OpenAI backend, added in an earlier cycle, had never been documented in `standalone-ai-guide.md` until this pass. The report explicitly states Fireworks AI and Cerebras were deliberately excluded, not overlooked.

The Security Assessment Toolkit (the platform's port-scanning/active-reconnaissance feature) went through two dedicated QA passes, `HORIZON_GRID_PORT_SCANNING_QA_REPORT.md` and `SECURITY_ASSESSMENT_QA_REPORT.md` , whose fixes are independently confirmed present in `backend/app/core/security_assessment.py` and `backend/app/security_assessment/nmap_tool.py` as read directly this turn. The most significant items:

Bug	Impact	Fix / Status
No way to cancel a running scan (no endpoint, button, or subprocess-kill path)	A slow or mistaken scan could only be waited out; a backend restart would orphan the running <code>nmap</code> process	Added <code>CANCELLED</code> status, cancel endpoint, task-tracking dict, dedicated <code>asyncio.CancelledError</code> handling in both the orchestrator and the nmap tool. Fixed and independently confirmed in code (<code>security_assessment.py</code> lines 91-165, <code>nmap_tool.py</code> lines 146-155).
Migration added lowercase <code>'cancelled'</code> against Postgres's existing uppercase enum labels	2 tests failed with <code>invalid input value for enum ... "CANCELLED"</code>	Changed to <code>ADD VALUE IF NOT EXISTS 'CANCELLED'</code> . Fixed.
Two concurrent cancel requests for the same run could each write their own audit row	Audit trail double-logged a single cancellation	The live-task path's own background task is now the sole writer; the restart-orphan path gates its write on <code>rowcount</code> . Fixed and confirmed in code (<code>security_assessment.py</code> lines 123-137).
IPv6 scans never actually probed the target — nmap needs an explicit <code>-6</code> flag for an IPv6 literal, otherwise it exits 0 having scanned 0 hosts, indistinguishable from a clean scan	High — IPv6 scanning was silently non-functional	Added <code>-6</code> conditionally. Fixed and confirmed present in code (<code>nmap_tool.py</code> lines 124-131).
An unknown or mistyped scan profile id was silently accepted and reported as a successful, no-findings <code>completed</code> scan	High — indistinguishable from a real clean scan	Validate the profile id up front, return <code>400</code> . Fixed and confirmed in code (<code>security_assessment.py</code> lines 298-313).
A malformed CIDR target crashed the run endpoint with a raw <code>500</code>	Medium	Catch the <code>ValueError</code> , return <code>400</code> . Fixed and confirmed in code (lines 260-266).
A run's own successful completion could silently overwrite a status another writer (e.g. startup orphan-recovery) had already finalized	Medium — self-contradictory run rows	Guard the completion update with <code>WHERE status IN (PENDING, RUNNING)</code> . Fixed and confirmed in code (lines 462-474).
A test-isolation race: <code>status</code> flips to <code>COMPLETED</code> before the post-completion AI-refresh step finishes, letting a test tear down its DB engine mid-refresh	Flaky/corrupting test suite	Added <code>wait_for_background_runs()</code> , a real synchronization primitive. Fixed and confirmed present in code (lines 168-177).

A handful of additional items in these two reports — empty `tool_ids` accepted as a no-op run, a Spamhaus DNS-rejection misread as "malicious," a confusing generic 401 on a last-admin-protection race, an AI assessment citing findings while its structured evidence list stayed empty, and a CI-only mocking gap for `is_available()` — are reported as fixed but were not independently re-verified in source this turn, since they fall outside the two files (`security_assessment.py`, `nmap_tool.py`) actually read. Both reports also disclose limitations that were not treated as bugs: no live elevated Windows installer run or Linux `.deb` /`systemd` run was performed for either toolkit version at the time; TLS-downgrade testing and DNS zone-transfer/subdomain enumeration are out of scope by design; `local_observation` / `ai_interpretation` provenance categories exist in the schema but are unpopulated; and cross-run correlation merging does not re-run full cross-provider dedup across old and new results.

The deterministic scoring engine (`backend/app/scoring/engine.py`, version-stamped `SCORING_ENGINE_VERSION = "1.0"`) has its own dedicated unit-test file, `backend/app/tests/unit/test_scoring_engine.py`, read directly this turn. It exercises: zero-evidence handling, non-OK provider results being ignored, corroboration scaling with agreeing-provider count, confidence collapsing (not averaging) on conflicting verdicts, qualifying vs. non-qualifying correlation edges, a documented single-provider-flood regression (a fix for a confirmed exploit where one account listing several distinct malware-family tags across free-text fields could otherwise saturate the correlation component alone), the Security Assessment severity floor not inflating `malicious_probability`, severity-band thresholds, and the `engine_version` stamp on every result. A separate integration test, `test_deterministic_scoring.py`, runs a real end-to-end lookup with a stubbed AI client deliberately returning invented risk numbers and asserts the persisted score matches the deterministic engine, not the AI's invented values, while the AI's verdict *choice* (e.g. "suspicious") is preserved. This integration test is gated by a `pytest.mark.skipif` on live Postgres/Redis reachability; it was not run in this session, so only what it asserts — not its current pass/fail status — can be reported here.

Adversarial Testing: A Live Red-Team Pass Against the Running System

The single largest QA document in the project's history is `ENTERPRISE_QA_PHASE1_FINDINGS.md`, a 1241-line, 21-dimension adversarial audit run directly against the live, running stack — authentication and session lifecycle, RBAC and privilege escalation, multi-admin concurrency races, last-admin-protection guards, cross-user data isolation, AI-provider credential lifecycle, AI-backend chaos/outage simulation, IOC-provider fan-out resilience, AI hallucination and evidence-fidelity defenses, real critical-IOC classification runs, the Security Assessment Toolkit's safety boundaries (including live nmap scans against `127.0.0.1`), command-injection probing, file-export/path-traversal surface, real-browser frontend chaos testing via Puppeteer, audit-log integrity, secret scanning, dependency/CVE inventory cross-checked against live OSV.dev queries, backend observability, and fresh-eyes documentation validation. It is explicitly an audit/red-team report, not a fix report: it documents 74 real findings numbered against the running system, none of them fixed within the document itself (plus two literally corrupted placeholder entries in the compiled output, flagged in the report itself as a data artifact rather than a finding). A subsequent adversarial-verification pass re-confirmed the 7 most severe (P1) findings, leaving the rest reported but not independently re-verified a second time.

The 7 confirmed P1 findings:

Finding	Root cause (as documented)
An oversized IOC value can crash the <code>/lookup/stream</code> SSE generator and leave the investigation permanently stuck at <code>status='running'</code>	OTX's <code>source_url</code> has no length cap and can exceed the <code>provider_results.source_url</code> <code>VARCHAR(2048)</code> column; the generator's own exception handler then attempts a second <code>commit()</code> on an already-rolled-back session, raising an uncaught <code>PendingRollbackError</code> that escapes the generator before its "mark FAILED" safety net can run
No request-correlation ID anywhere in the backend	<code>app/main.py</code> registers only <code>CORSMiddleware</code> and a Prometheus instrumentator; concurrent operations are unattributable from <code>docker logs</code> beyond spotting literal IOC-value substrings
Failed logins are invisible in stdout logs beyond a bare <code>401</code>	<code>login()</code> in <code>app/api/routes/auth.py</code> has zero <code>logger.*</code> calls; the attempted email reaches only a free-text audit field, with <code>actor_email</code> left NULL
Security Assessment Toolkit runs, including real nmap scans, produce no log output at all on success	No tool adapter has any <code>logger.*</code> call; the service layer only logs on failure
<code>ADMIN_GUIDE.md</code> self-contradicts on whether the provider/audit-log UI exists	Both features are real and fully working, live-verified against <code>GET /api/v1/runtime/ai-providers</code> (200 for admin, 403 for analyst)
<code>USER_GUIDE.md</code> claims "exactly eight routes... nothing else exists" and "no admin/user-management UI... no audit log"	Both <code>/admin</code> and <code>/providers</code> are real, reachable, RBAC-gated pages
<code>USER_GUIDE.md</code> 's breakdown of <code>/lookup/new</code> omits the entire Security Assessment feature and the AI Comparison section	None of the three core docs mention Security Assessment at all

Below that severity tier, the report documents a long list of P2-P4 findings that were reported with live evidence but not independently re-verified in the adversarial pass: a ~26x login-timing side channel that leaks account existence (bcrypt only runs for real emails: median 231.4ms for an existing-email/wrong-password attempt vs. 8.8ms for a nonexistent email, N=40/arm); no rate limiting on

`/auth/login` (15 rapid failed logins all returned 401 with no throttling); an audit-log write on every `set_user_active()` call even when nothing changed, producing 3 audit rows for 1 real transition under concurrency; an admin-console row-lock (`_lock_all_admin_rows()`) taken unconditionally even for an unrelated non-admin edit, measured to block an unrelated PATCH for ~2.5 seconds; a final-assessment retry limited to exactly one attempt on `ValidationError` only, whose generic fallback text is indistinguishable in the UI from a true AI-backend outage; a `GET /providers/health` configured field that reads a stale, import-time class attribute rather than the actual runtime-DB-backed credential state; CISA KEV's most authoritative fact for Log4Shell being silently dropped from a final assessment when the AI's per-provider summary generation failed on malformed JSON, with no deterministic fallback reconstructing it from the still-valid raw provider fields; a small local model (`gemma2:2b`) overclaiming certainty (100.0 probability / "high" confidence) against underlying evidence carrying only 0.5/"medium" confidence, with no cross-field validator catching it; a `ProviderSummary` object that can be internally self-contradictory (`reputation: 'unknown'` beside `threat_level: 'high'`) because only `FinalAssessment`, not `ProviderSummary`, has a schema validator; `agreeing_providers` / `disagreeing_providers` coming back empty in all 3 fresh test runs despite being documented as mandatory, breaking a UI citation feature; a Security Assessment Toolkit silent-failure path where a long-but-schema-valid target string can be committed as a `PENDING` run and then never run, complete, or fail, because an uncapped audit-log f-string overflows a `VARCHAR(1000)` column after the run row is already committed; unescaped CSV export enabling formula/CSV injection (CWE-1236); unescaped PDF export allowing a malformed ReportLab tag to crash the export or a well-formed tag to inject spoofed formatting into an "official-looking" report; `export_lookup()` gated on the broader `lookup:read` instead of the dedicated `lookup:export` permission, letting a VIEWER download real exports; a page refresh mid-investigation silently starting a duplicate lookup rather than resuming the original; an unhandled `IntegrityError` on concurrent "Add to Basket" for the same new IOC value; zero duplicate protection on "Add to Case"; cross-tab logout not propagating, leaving a second tab signed in indefinitely; weak, hardcoded default database credentials (`ioc:ioc@postgres`) confirmed live as the actual running Postgres credential; a hardcoded `jwt_secret_key` fallback with no startup guard if ever left unset (confirmed not the live value, since the wizard generates a real secret, but with no code-level safety net for a deployment path that skips the wizard); `next==14.2.15` carrying 30 OSV advisories including a critical middleware auth-bypass CVE, with this app's own specific attack surfaces confirmed unused but App Router RSC-cache advisories left unruled-out; `cryptography==43.0.1` carrying 11 OSV advisories relevant to this app's TLS tool and Fernet credential encryption; and a `pip-audit` false negative, where both the default and OSV-backed modes reported "no known vulnerabilities" for packages independently confirmed vulnerable via direct OSV.dev queries from the same container.

Lower-severity, live-observed findings include: 3 of 5 simultaneous investigations taking 87-160 seconds under concurrent load against a documented ~20-40 second single-run baseline (confounded by an unrelated concurrent QA process hammering the same endpoint, flagged as needing isolated re-measurement); a stale "last test: failed" badge persisting on a currently healthy AI backend; an internally incoherent risk score for `8.8.8.8` driven by five coincidentally matched, irrelevant OSINT hits; a confident "benign" verdict for the reserved TEST-NET-3 address `203.0.113.77` with zero real evidence, alongside a raw JSON-repair artifact leaking into a structured `recommended_actions` field; empty `mitre_mappings` for all 5 test IOCs including a textbook Log4Shell/T1190 case; an nmap argument-injection gap in `nmap_tool.py`'s target parameter, confirmed not currently exploitable because the attacker-controlled string occupies only one, always-terminal argv slot; an `ioc_type_hint` field letting a caller assign an arbitrary IOC type label bypassing format-regex detection; `structlog` configured in `main.py` but never actually used anywhere else in the codebase; high-frequency frontend polling (~every 250-300ms) of an AI-active-backend endpoint flooding `docker logs`; `/providers` rendering a blank content area with no access-denied message for a non-admin analyst whose backing APIs correctly 403; and a test-infrastructure bug in a shared `_delete_user()` teardown helper lacking foreign-key cleanup, causing 2 deterministic integration-test failures and leaking 9 orphaned test users, found and cleaned up but not code-fixed within the report.

Mission-Critical and Chaos Testing

An earlier, separate effort — `LAN_DEPLOYMENT_QA_REPORT.md` and `LAN_REGRESSION_TEST_REPORT.md`, dated 2026-08-12/2026-08-14, with no version number stated in either file — converted the platform from localhost-only to LAN accessibility. It is documented as distinct prior work, not part of the later v0.2.3 certification. Its one clearly attributable fix, per `LAN_DEPLOYMENT_QA_REPORT.md` line 20, is a `docker-compose.yml` fallback bug: `_${PUBLIC_API_URL}:-http://localhost:8000` (colon-hyphen) falls back to the hardcoded default even when the variable is explicitly set to empty, changed to `_${PUBLIC_API_URL}-`.

The dedicated mission-critical reliability effort is `MISSION_CRITICAL_CERTIFICATION_REPORT.md`, which states explicitly: "**Versión certified:** 0.2.3." Its 32-row requirements table and accompanying sections describe real chaos and failure-injection testing, not just functional checks. The Linux `.deb` installer was rebuilt and installed as a real upgrade over a live existing v0.2.1 instance running inside WSL, exercising the reconfigure wizard end-to-end: prerequisite gate, pre-upgrade backup, port auto-remap, new `.env`, containers reaching healthy, systemd units enabled. The Windows `.exe` was compiled and its build log inspected to confirm new scripts were packaged, but a live elevated Windows install was explicitly not performed, disclosed as blocked by a UAC prompt the autonomous session couldn't satisfy.

Chaos-style tests specifically confirmed live: `restart: unless-stopped` on all 8 Docker services was validated against a real OOM-killed container; a `GET /health/detailed` endpoint distinguishing `healthy` / `degraded` (Redis down, Postgres up) / `down` (HTTP 503) was validated by actually stopping and restarting Postgres; database backup/restore was validated end-to-end by backing up, inserting a marker row, restoring, and confirming the marker was gone and the platform healthy afterward; and a 3-hour, 171-cycle soak test against the live stack showed flat memory throughout (backend 114.6→106.8 MB, Postgres 37.98→41.67 MB, Redis 10.23→9.03 MB) with zero spontaneous container restarts.

Real defects found and fixed during this certification, with status as stated in the report:

- The Linux `horizon-grid.service` systemd unit declared `WantedBy=multi-user.target` but `systemctl enable` was never actually called anywhere in the codebase — only `daemon-reload` ran in `postinst` — meaning a working install would silently not survive a host reboot. Self-discovered during this pass, fixed by adding the enable call to the wizard's success path, verified live via `systemctl is-enabled` returning `enabled`.
- `check-prerequisites.sh`'s JSON output had a permanently empty `checks` array, because bash's lowercase `true` / `false` interpolated as bare Python identifiers caused a silent `NameError` on every call, swallowed by an error fallback; console output and exit code were computed separately and unaffected, which is why it had gone unnoticed. Fixed by passing fields through the environment instead of interpolating into generated source, verified live for both pass and simulated-failure cases.
- `BaseProvider.run()` was swallowing exception types the orchestrator's retry loop needed to see, meaning `provider_max_retries` had no effect for providers without their own network-error handling. Fixed and live-verified with a fake provider showing a `ConnectError` genuinely retried.
- The Final Assessment panel showed an identical badge for a genuine AI failure and a correct no-evidence skip, because the backend already tracked and sent an `ai_outcome` field the frontend never read. Fixed and confirmed against a real `ai_outcome="failed"` record.
- A full-suite regression run had reported 365 passed, matching the previous known-good count, traced to a test container (`hgmc`) running a stale image predating the session's own backend changes. The image was rebuilt and the suite re-run for real, reporting 380 passed, 39 skipped.
- Redis is stated to now persist across restarts, described as previously not persisting, confirmed live.

The report is equally explicit about what remained open at the end of this pass, per its own §9: no automated host-disk-space alerting; no retention/cleanup job for growing investigation tables; Neo4j and OpenSearch fully provisioned with zero actual read/write traffic, flagged as an open resource-cost question rather than fixed; a DNS-rebinding TOCTOU gap on the SSRF check (it validates a DNS snapshot, while the real outbound call re-resolves independently); no off-site or off-host backup copy (local-disk-only); no global ONLINE/LIMITED/OFFLINE UI indicator; no real simulated power-loss (mid-write) test performed; per-subprocess (nmap) OS-level limits remaining container-level only; `structlog` still configured but unused anywhere in application code; the full ~40-document existing doc set not exhaustively regenerated; and, again, no live elevated Windows installer run performed in that session.

Installer Testing

Installer-level testing recurs across nearly every phase of this project's history rather than being a single discrete event. The earliest installer bugs were the Setup Wizard's session-token and `.env`-escaping issues covered above, from `API_CONFIGURATION_FIX_REPORT.md`. The v0.2.0 AI-expansion pass found and fixed the Windows installer bundling the entire local dev virtualenv (~190 MB of unused content), confirmed via `dpkg-deb -c`-style content inspection for the analogous Linux case in later work. The v0.2.3 Mission-Critical Certification actually installed the rebuilt Linux `.deb` as a

real upgrade over a live v0.2.1 instance in WSL and ran its reconfigure wizard end-to-end, while explicitly disclosing that no live elevated Windows install was performed in that session, for the UAC-related reason stated above.

The most recent installer-specific findings come from `FULL_FUNCTIONAL_TEST_REPORT.md` (dated 2026-08-21, version 0.2.3), which found and fixed two P0-severity installer defects: a stale shipped `.exe` that was missing `docker-compose.yml` entirely, and a leftover registry `InstallLocation` value from an old test harness that silently redirected new installs to a fake path. It also found and fixed one P1 AI-schema defect: a free-text `threat_assessment` field able to state a conclusion contradicting the validated `final_verdict` field. That same report discloses that the Setup Wizard's own WinForms GUI was never clicked through end-to-end in this pass — equivalent PowerShell functions were invoked directly instead — and that a Windows-reboot phase was deliberately deferred because the user was asleep at the time, with Linux install-plus-reboot and uninstall/reinstall phases also not yet run.

The earlier 0.1.0-era release cycle documents a separate installer verification step in `FINAL_SHIP_AUDIT_REPORT.md` : a release-candidate lock/audit that recomputed and verified the exact shipped installer's SHA256 hash byte-for-byte against the fixes it was meant to contain, then ran a final smoke test of the full user journey. A same-day addendum to that report describes a follow-up fix pass (implementing real PDF/CSV export and closing a concurrency edge case) that produced a new installer build, with a new SHA256 hash (`1FD8349E6E2027B51F72485D74AB8EF6B1EF0281C74F86DEB2E66226CC6FCACD`) and 180 tests passing, explicitly stating that this addendum "does not change the final verdict above."

Final Release QA: A History of Verdicts

The final-release phase of QA is not one pass with one verdict — it is a sequence of dated reports, each stating its own verdict for the specific product/build state it describes, and none of the source files reconcile all of them into a single authoritative statement. In chronological order, as actually documented:

`FINAL_QA_REPORT.md` , dated 2026-08-14, against the already-installed IOC Intelligence Platform, is a red-team regression pass that found and live-fixed three real bugs — a credential-wipe-on-save bug, a silent AI-assessment failure on a self-contradictory verdict, and unbounded provider data burying severity signal in the AI prompt — while also flagging the export feature as apparently completely non-functional and finding live plaintext secrets in a repo-root `.env` . Its stated verdict, line 27: **"NOT RELEASE READY,"** on the grounds that the credential-wipe bug, the non-functional export claim, and the plaintext-secrets exposure were each independently disqualifying per that report's own severity rules.

`RELEASE_BLOCKER_TRACKER.md` tracked those same bugs through fix-test cycles. BUG-01 needed a second round to add row-locking and real database-backed regression tests after an adversarial re-check found gaps; BUG-03 needed two further correction rounds after its own fix caused new regressions — first breaking MITRE ATT&CK's long descriptions, then breaking Groq's request-size limit. BUG-04, the apparent export non-functionality, was retracted entirely: the original test had used the wrong reproduction method, and JSON/Markdown export genuinely worked client-side all along; only PDF/CSV correctly 404'd, which was already-documented, non-regression behavior at that time.

`FINAL_SHIP_AUDIT_REPORT.md` , the same day, 2026-08-14, same product and version (0.1.0), is a release-candidate lock/audit rather than a fresh QA pass. Its stated verdict, line 157: **"RELEASE READY WITH KNOWN LIMITATIONS."** Its disclosed known limitations at that point were PDF/CSV export not implemented, a narrow first-time-configure concurrency edge case, AI verdict quality depending on the chosen model's capability, and Groq's shared-tier rate limit as a real, correctly-handled external constraint. Its same-day addendum then closed the first two of those (implementing real PDF/CSV export and fixing the concurrency edge case) without changing the stated verdict.

`KNOWN_LIMITATIONS.md` , titled for "IOC Intelligence Platform 0.1.0," is the resulting limitations ledger for that release, reflecting the post-addendum state (PDF/CSV listed as implemented, 180 automated tests referenced).

`FINAL_RELEASE_QA_REPORT.md` documents a separate, later effort under the "HORIZON GRID" rebrand — new providers, a deterministic scoring engine, a dashboard, and an installer password-mismatch fix — after which several real bugs found mid-mission were fixed and live-verified, including a user-confirmed real installer run. Its stated verdict, line 133: **"VERDICT: RELEASE READY WITH KNOWN LIMITATIONS."** Its own §13 discloses that the AI-generated executive summary can be slow under heavy concurrent load with a local Ollama backend (an infrastructure characteristic, not a bug); that `_provider_votes()` 's NaN/Infinity handling is a disclosed latent hardening gap with no live exploitable path found; and that the "Test Connection" button never falls back to an already-saved credential, a deliberate security property that is nonetheless a real point of user confusion.

`FULL_FUNCTIONAL_TEST_REPORT.md` , dated 2026-08-21 — the newest of these files by timestamp, and covering version 0.2.3 — is marked at the very top, line 3, **"Status: IN PROGRESS."** It found and fixed the two P0 installer bugs and the one P1 AI-schema bug described in the Installer Testing section above. Its own stated verdict, line 374: **"FUNCTIONALLY VERIFIED WITH DOCUMENTED LIMITATIONS,"** and the report itself immediately qualifies this, line 375, as **"an interim, not final, verdict,"** stating that roughly a third of a planned 33-phase mission was covered and that the remaining phases "are explicitly NOT YET DONE, not silently assumed to pass." Phases explicitly not covered in this pass, per lines 336-370: log rotation (Phase 21), performance under load (Phase 22), a broader security/RBAC sweep (Phase 23), a Windows reboot test deliberately deferred because the user was asleep (Phase 24), Linux install-plus-reboot (Phase 25), uninstall/reinstall (Phase 26), full regression and a soak test (Phases 27-28), and a final post-fix repeat of the critical path (Phase 31).

The Final, Most Recent Release Verdict

The most recent report by date, `FULL_FUNCTIONAL_TEST_REPORT.md` (2026-08-21, v0.2.3), explicitly declines to state a final release verdict. Its own top-of-document status is **"IN PROGRESS,"** and its stated verdict — **"FUNCTIONALLY VERIFIED WITH DOCUMENTED LIMITATIONS"** — is qualified in the same document as **"an interim, not final, verdict."** No source file available for this review reconciles that report with the separately dated `MISSION_CRITICAL_CERTIFICATION_REPORT.md` (also version 0.2.3, whose own stated verdict, line 109, is **"MISSION-CRITICAL READY WITH DOCUMENTED LIMITATIONS"**) into one single authoritative statement covering the whole product. Each report's verdict applies only to the scope and build state described within that same report, and the most recent one on record says plainly that the overall picture is still incomplete. Stated exactly, without upgrading or softening either: as of the newest dated report, the project's own most current self-assessment is an explicitly interim **"FUNCTIONALLY VERIFIED WITH DOCUMENTED LIMITATIONS,"** alongside a separately scoped, non-interim **"MISSION-CRITICAL READY WITH DOCUMENTED LIMITATIONS"** for the reliability-hardening certification specifically — neither an unqualified "release ready" nor a "not release ready" verdict exists in the most current material.

Known Limitations

The following are limitations the source reports themselves disclose, organized by the report that states them — not inferred, extrapolated, or reworded into a stronger or weaker claim than the source states.

From `KNOWN_LIMITATIONS.md` (IOC Intelligence Platform 0.1.0): Groq's free/shared-tier rate limit (12,000 tokens/minute) is a real external constraint, handled cleanly (one failure, not a retry storm) but not something the platform controls. AI-generated verdicts are model-dependent in quality, most visibly with smaller local models such as Ollama's default `llama3.2:3b` . No literal GUI-driven installer check-through was performed for that release's build, blocked by the test environment rather than the product. As user action required: remove or reset the disposable QA test account (`final-rtp-admin@examp1e.com`) and test case ("Final Ship Audit Smoke Test Case") before handing the environment to a real end user. As a security action required: rotate all exposed API credentials, including an AWS key pair and several threat-intel provider keys, found in a dev-environment `.env` file — a pre-existing exposure, not introduced by the release, and not rotated by the QA process itself.

From `FINAL_SHIP_AUDIT_REPORT.md` : AI verdict quality depends on the chosen model's capability. Groq's shared-tier rate limit is a real external constraint, correctly handled. (Its other two pre-addendum limitations — PDF/CSV export not implemented, and a narrow first-time-configure concurrency edge case — were closed by the same-day addendum and are no longer current limitations.)

From `FINAL_RELEASE_QA_REPORT.md` §13: The AI-generated executive summary can be slow under heavy concurrent load when a local Ollama backend is configured — an infrastructure characteristic, not a defect. `_provider_votes()` 's NaN/Infinity handling is a disclosed latent hardening gap with no live exploitable path found today. The "Test Connection" button never falls back to an already-saved credential — a deliberate security property, not a bug, but a real source of user confusion.

From `MISSION_CRITICAL_CERTIFICATION_REPORT.md` **§9 (v0.2.3)**: No automated host-disk-space alerting. No retention/cleanup job for growing investigation tables. Neo4j and OpenSearch are fully provisioned with zero actual read/write traffic, an open resource-cost question rather than a fixed issue. A DNS-rebinding TOCTOU gap remains on the SSRF check (it validates a DNS snapshot; the real outbound call re-resolves independently). No off-site or off-host backup copy exists (local-disk-only). No global ONLINE/LIMITED/OFFLINE UI indicator exists. No real simulated power-loss (mid-write) test was performed. Per-subprocess (nmap) OS-level limits remain container-level only. `structlog` is configured but unused anywhere in application code. The full ~40-document existing doc set was not exhaustively regenerated. No live elevated Windows installer run was performed in that session.

From `FULL_FUNCTIONAL_TEST_REPORT.md` **(v0.2.3, 2026-08-21)**: Log rotation, performance under load, a broader security/RBAC sweep, a Windows reboot test, Linux install-plus-reboot, uninstall/reinstall, full regression, a soak test, and a final post-fix repeat of the critical path are all explicitly not yet covered by this pass, not silently assumed to pass. The Setup Wizard's own WinForms GUI was not clicked through end-to-end in this pass; equivalent PowerShell functions were invoked directly instead.

From `HORIZON_GRID_AI_PROVIDER_EXPANSION_QA_REPORT.md` : Fireworks AI and Cerebras are deliberately excluded from the AI-backend list, not an oversight.

From `HORIZON_GRID_PORT_SCANNING_QA_REPORT.md` **and** `SECURITY_ASSESSMENT_QA_REPORT.md` : No live elevated Windows installer run or live Linux `.deb` install/systemd run was performed for either toolkit version described in those reports. TLS downgrade testing and DNS subdomain enumeration/zone transfers are out of scope by design. `local_observation` / `ai_interpretation` provenance categories are modeled in the schema but unpopulated. Cross-run correlation merging does not re-run full cross-provider dedup/confidence-boost logic across old and new results. Vulnerability enrichment depends on nmap's own service/version banner detection quality.

Roadmap and Conclusion

What's Provisioned but Not Yet Wired Up

Two pieces of infrastructure already run in the Docker Compose stack with real configuration fields but no consuming code yet, confirmed directly against source: **Neo4j** (the `neo4j:5-community` container with the APOC plugin, plus `neo4j_uri` / `neo4j_user` / `neo4j_password` settings — no driver instantiation or Cypher code exists anywhere in the backend today; the correlation graph the product actually shows lives entirely in Postgres) and **OpenSearch** (the `opensearchproject/opensearch:2.17.0` container plus an `opensearch_url` setting — no indexing or query code exists against it today). Completing either is a bounded extension of infrastructure already provisioned, not a new subsystem. The full detail, including a diagram of today's actual wiring versus the provisioned-but-unconsumed pieces, lives in the Future Roadmap chapter of the Technical Documentation.

What's Genuinely Open, Disclosed by the Project's Own QA Reports

This submission's Bug History and Testing Journey chapters carry forward, rather than hide, what the project's own most recent QA reports state as still open at the time they were written: no automated host-disk-space alerting or investigation-table retention job; a DNS-rebinding TOCTOU gap on the outbound SSRF check; no off-site backup copy; several P2–P4 findings from the adversarial red-team pass not yet independently re-verified a second time; and the most recent functional-test report explicitly marking its own verdict as interim, with roughly two-thirds of a planned 33-phase mission not yet run. None of this is softened in the narrative chapters — the project's engineering discipline throughout has been to disclose a gap the moment it's found, not to wait until it's closed to mention it.

Where the Architecture Is Already Positioned to Grow

The registry pattern behind both the 18-provider IOC ecosystem and the 11-backend AI layer already normalizes heterogeneous sources behind one shared interface each (`BaseProvider` / `supports(ioc_type)`, and `call_claude_json()` respectively) — onboarding further providers or AI backends is proven, incremental work at this point, not a redesign. The same is true of the Security Assessment Toolkit's tool-adaptor pattern and the deterministic scoring engine's versioned weight table, both built to be extended rather than rewritten.

Conclusion

HORIZON GRID's central bet is that a threat-intelligence tool earns trust by showing its work: a deterministic score computed the same way every time, an AI that narrates a number it cannot override, a Provider Health page that reports Unknown rather than faking Healthy, and an Executive Dashboard where every tile is a real query result. That bet shaped the architecture from the scoring engine outward, and it shaped how this submission itself was written — every bug in the Bug History chapter is real, every open item in the Known Limitations sections is still open, and every screenshot in this package is a live capture of the running product, not a mockup.

The name and the tagline were chosen to describe exactly that outcome, not just to sound serious: **HORIZON GRID — Every Signal. One Operational Picture.**